

Profiling tools

Operating Systems



Outline

Performance and I/O

- Perf tool
- IO stat tool

Linux perf tools

- Install perf:

```
$ sudo apt-get install linux-tools-common linux-  
tools-4.2.0-27-generic linux-cloud-tools-4.2.0-27-generic
```

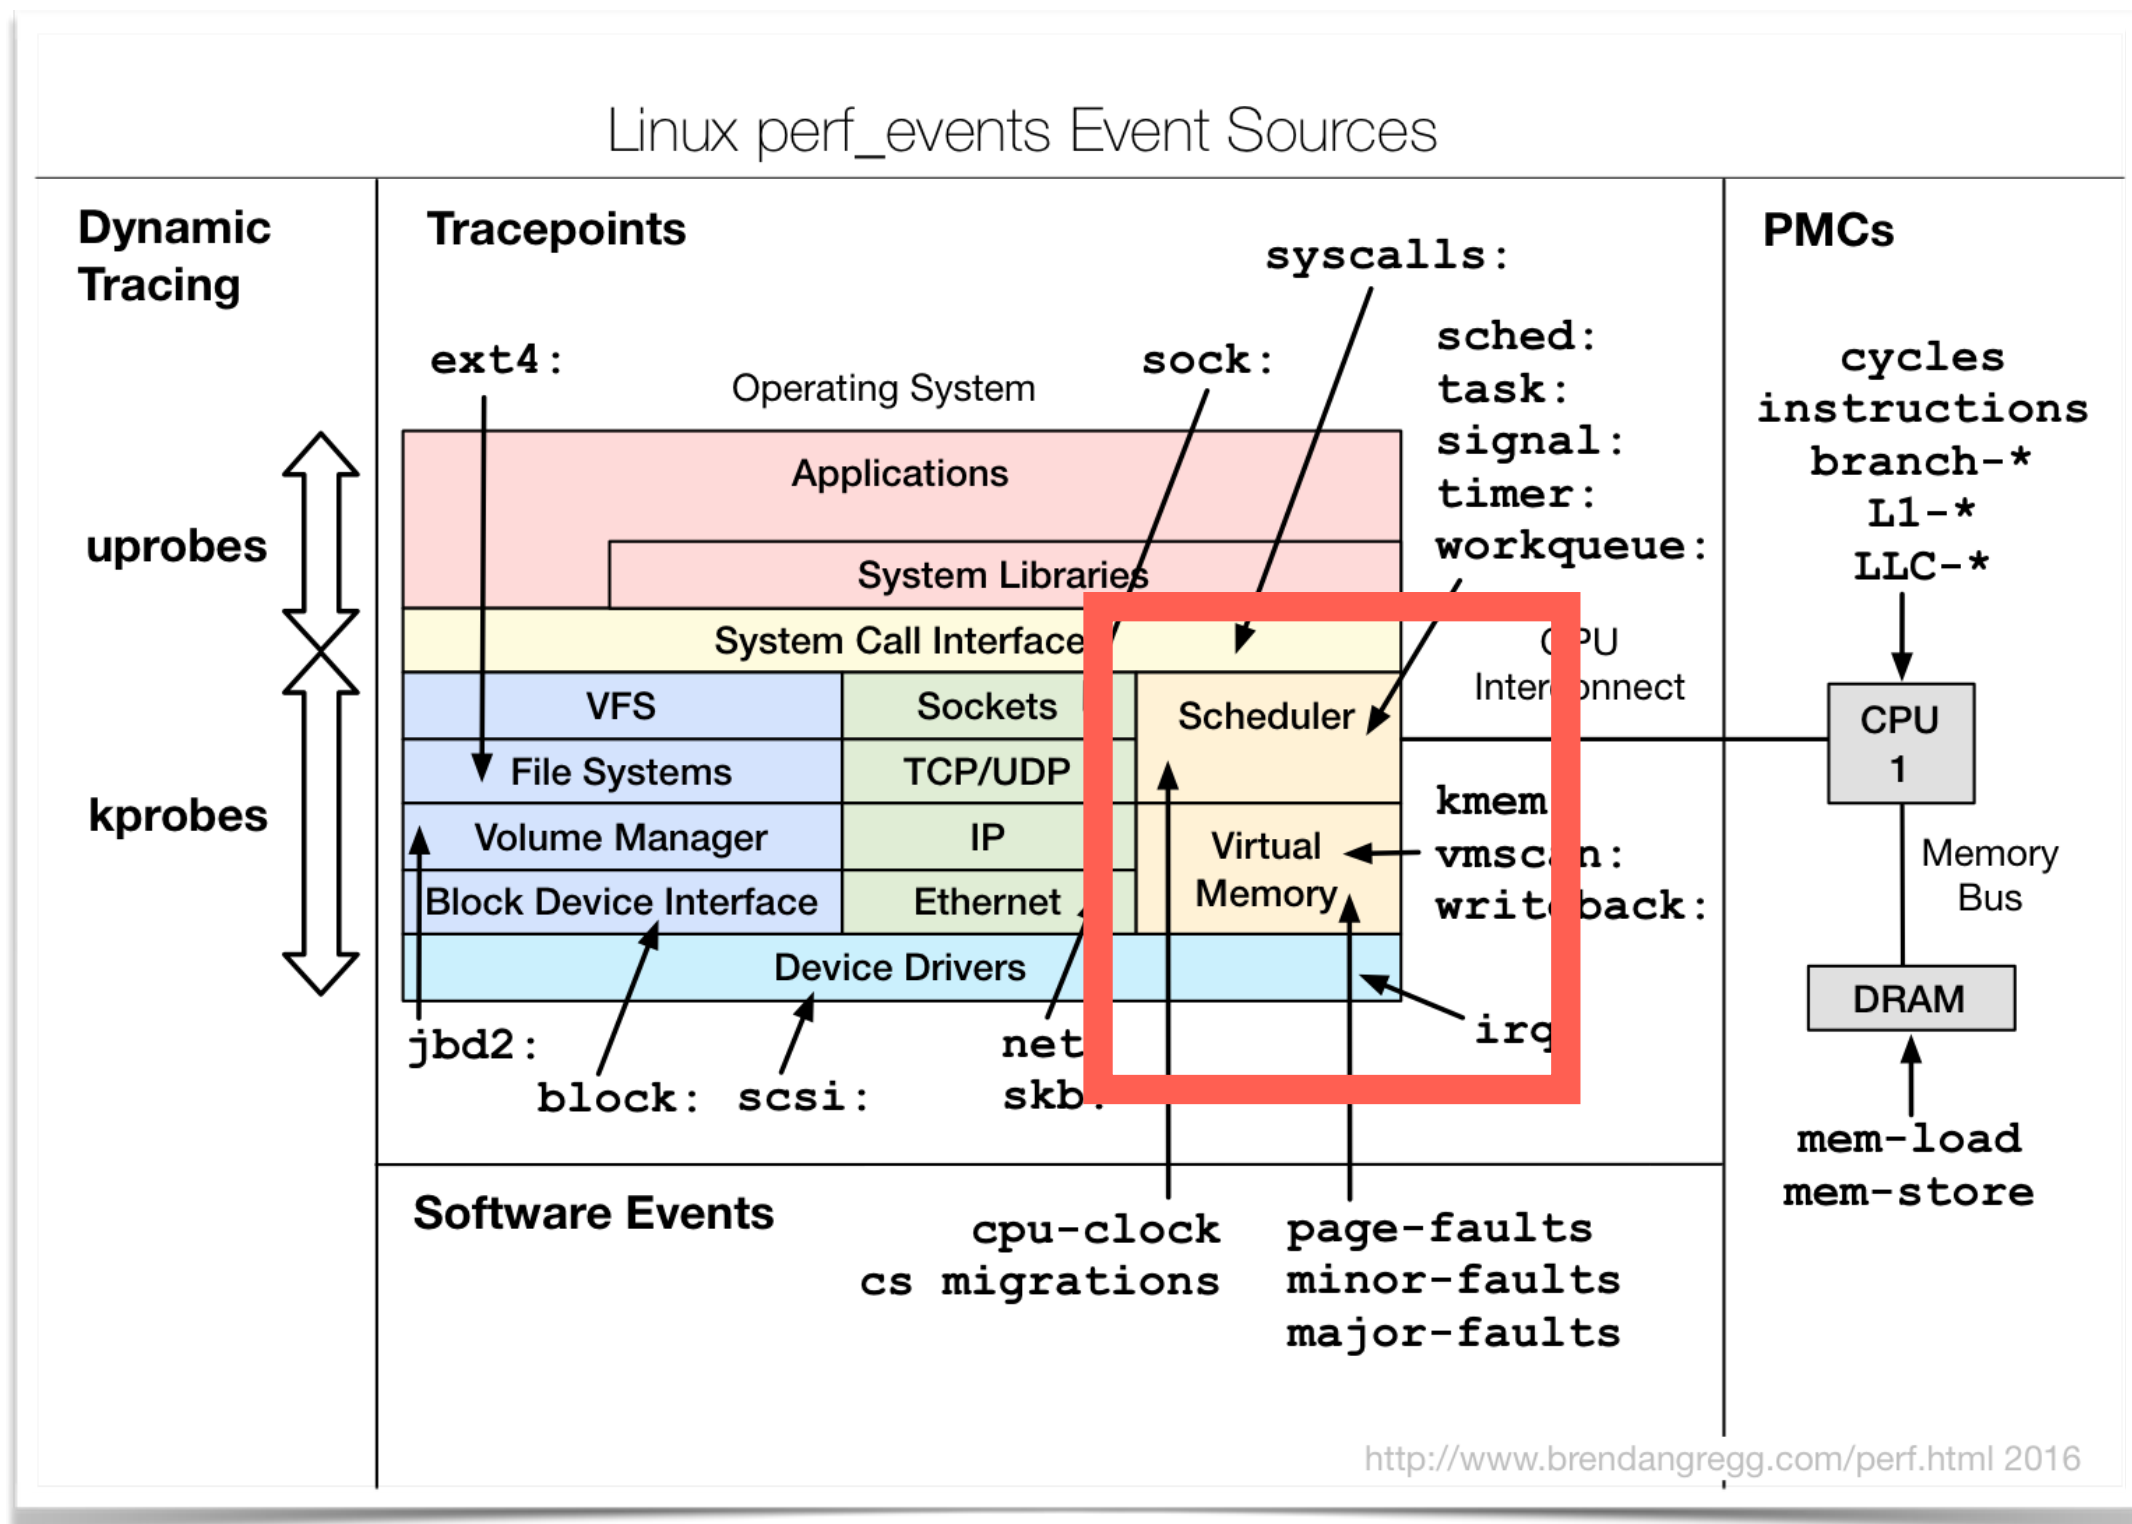
- Analyze cache and TLB:

```
$ perf stat -e cache-misses <command>  
$ perf stat -e dTLB-load-misses,iTLB-load-misses <command>
```

- Examples:

```
$ perf stat -e cache-misses ls>/dev/null  
$ perf stat -e dTLB-load-misses,iTLB-load-misses ls>/dev/  
null
```


Perf and virtual memory



Example: Perf counters

[HTTP://WWW.BRENDANGREGG.COM/PERF.HTML](http://www.brendangregg.com/perf.html)

CPU counter statistics for the specified command:

```
perf stat command
```

Detailed CPU counter statistics (includes extras) for the specified command:

```
perf stat -d command
```

CPU counter statistics for the specified PID, until Ctrl-C:

```
perf stat -p PID
```

CPU counter statistics for the entire system, for 5 seconds:

```
perf stat -a sleep 5
```

Various basic CPU statistics, system wide, for 10 seconds:

```
perf stat -e cycles,instructions,cache-references,cache-misses,bus-cycles -a sleep 10
```

Various CPU level 1 data cache statistics for the specified command:

```
perf stat -e L1-dcache-loads,L1-dcache-load-misses,L1-dcache-stores command
```

Various CPU data TLB statistics for the specified command:

```
perf stat -e dTLB-loads,dTLB-load-misses,dTLB-prefetch-misses command
```

Various CPU last level cache statistics for the specified command:

```
perf stat -e LLC-loads,LLC-load-misses,LLC-stores,LLC-prefetches command
```

Using raw PMC counters, eg, counting unhalted core cycles:

```
perf stat -e r003c -a sleep 5
```

Count all vmscan events, printing a report every second:

```
perf stat -e 'vmscan:*' -a -I 1000
```

Example: Perf counters

PMCs: counting cycles and frontend stalls via raw specification:

```
perf stat -e cycles -e cpu/event=0x0e,umask=0x01,inv,cmask=0x01/ -a sleep 5
```

Count syscalls per-second system-wide:

```
perf stat -e raw_syscalls:sys_enter -I 1000 -a
```

Count system calls by type for the specified PID, until Ctrl-C:

```
perf stat -e 'syscalls:sys_enter_*' -p PID
```

Count system calls by type for the entire system, for 5 seconds:

```
perf stat -e 'syscalls:sys_enter_*' -a sleep 5
```

Count scheduler events for the specified PID, until Ctrl-C:

```
perf stat -e 'sched:*' -p PID
```

Count scheduler events for the specified PID, for 10 seconds:

```
perf stat -e 'sched:*' -p PID sleep 10
```

Count ext4 events for the entire system, for 10 seconds:

```
perf stat -e 'ext4:*' -a sleep 10
```

Count block device I/O events for the entire system, for 10 seconds:

```
perf stat -e 'block:*' -a sleep 10
```

Recall: Memory Hierarchy


...

Recall: Memory Hierarchy

L1 Cache
~32KB, ~few ns


...

Recall: Memory Hierarchy

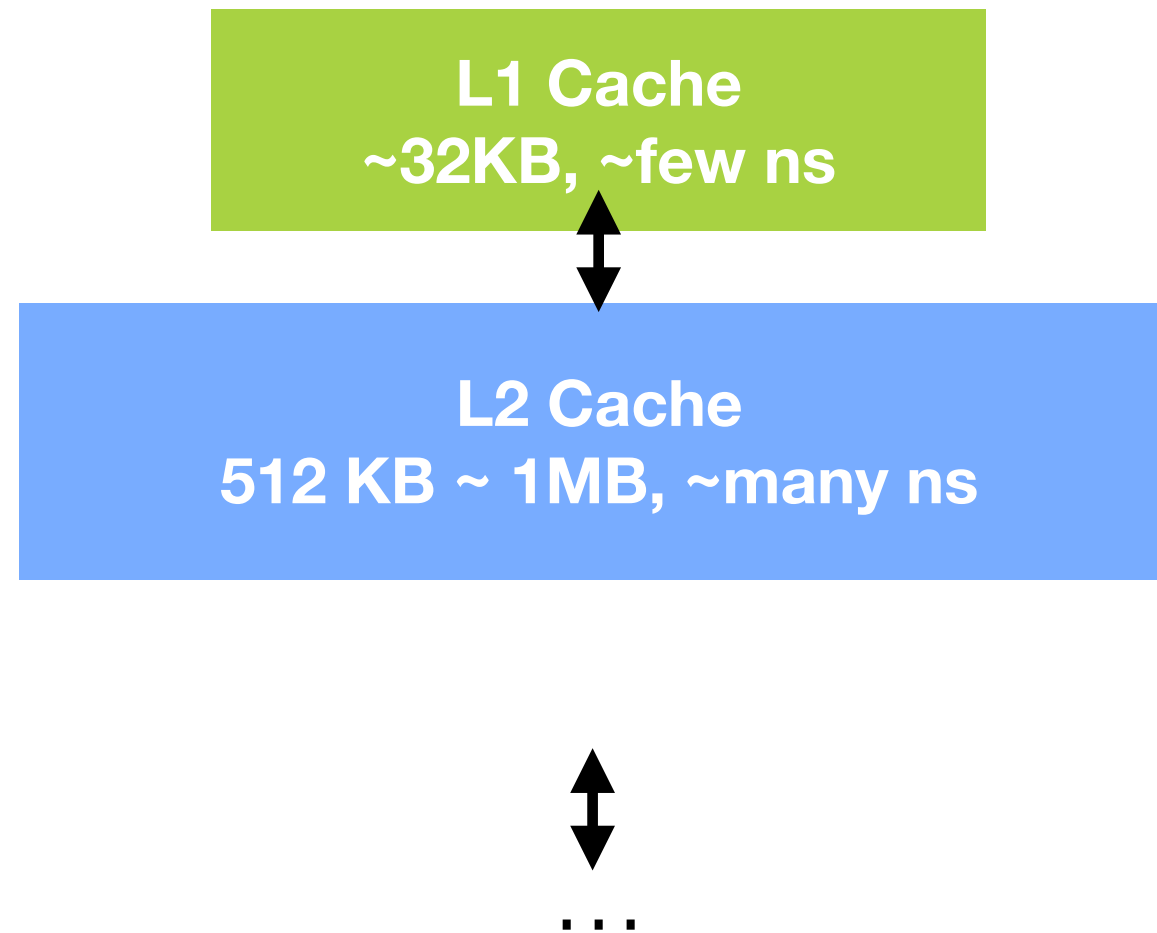


...

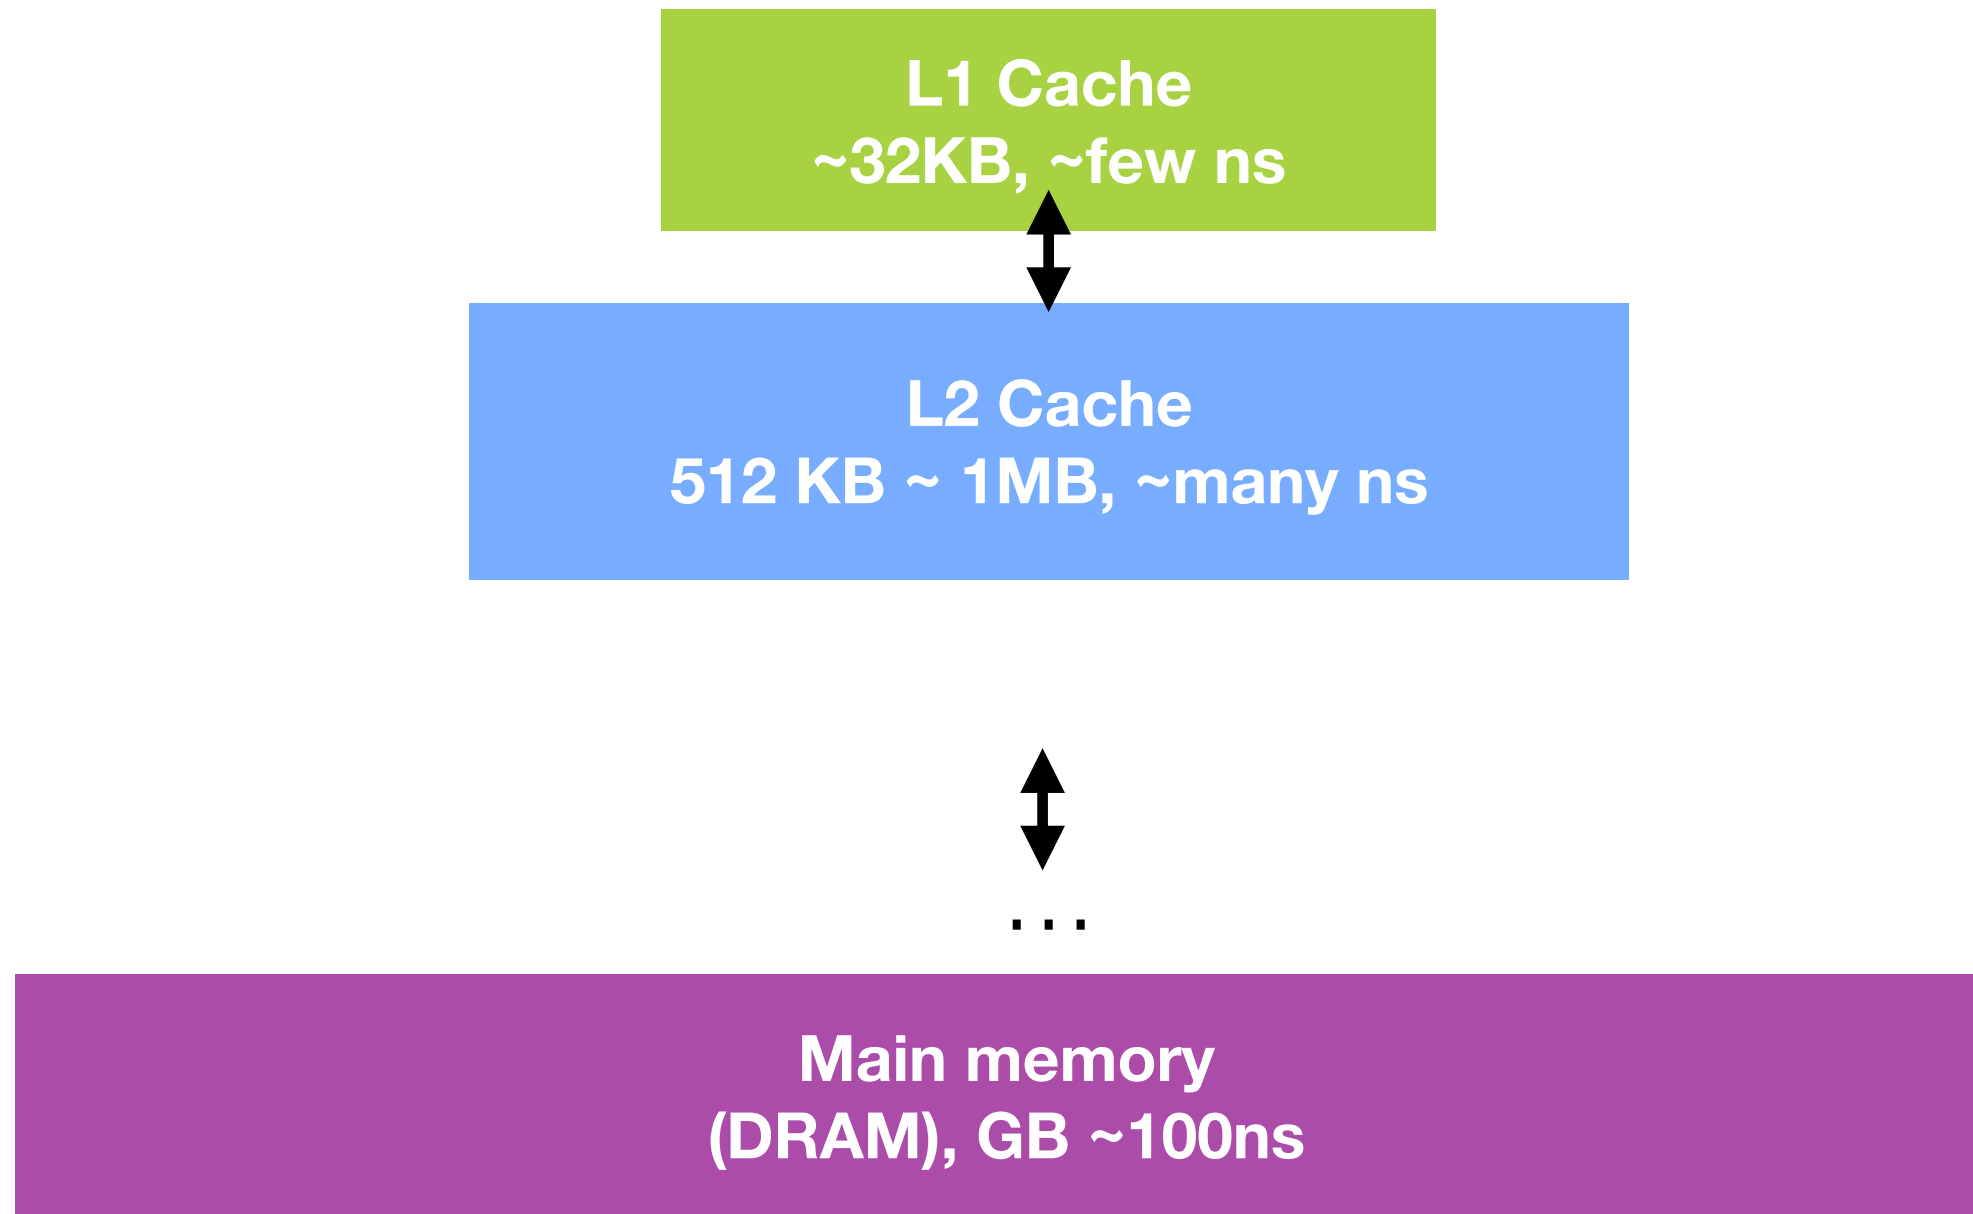
Recall: Memory Hierarchy



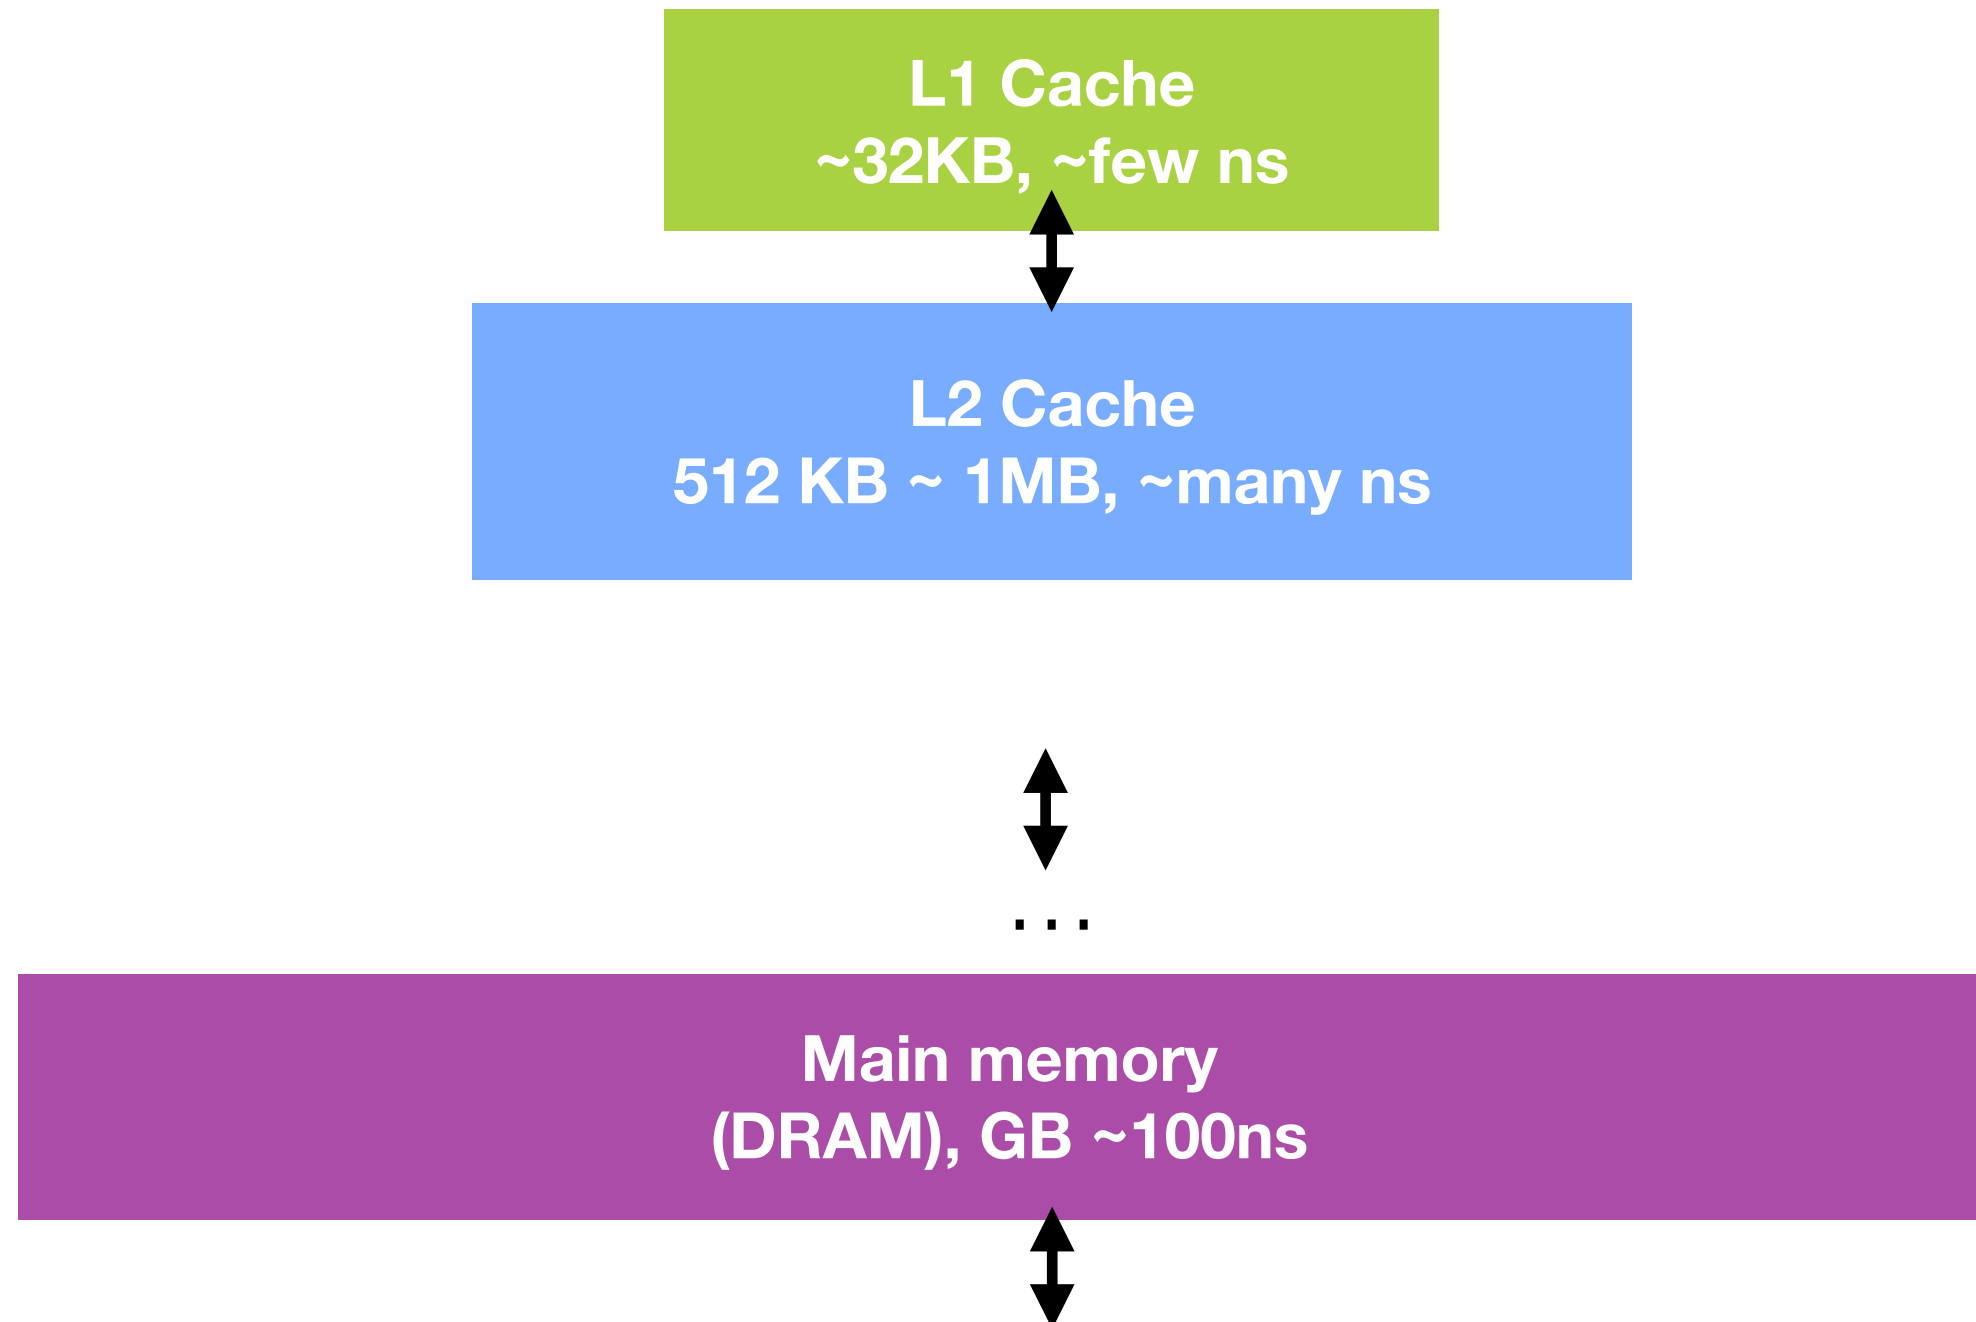
Recall: Memory Hierarchy



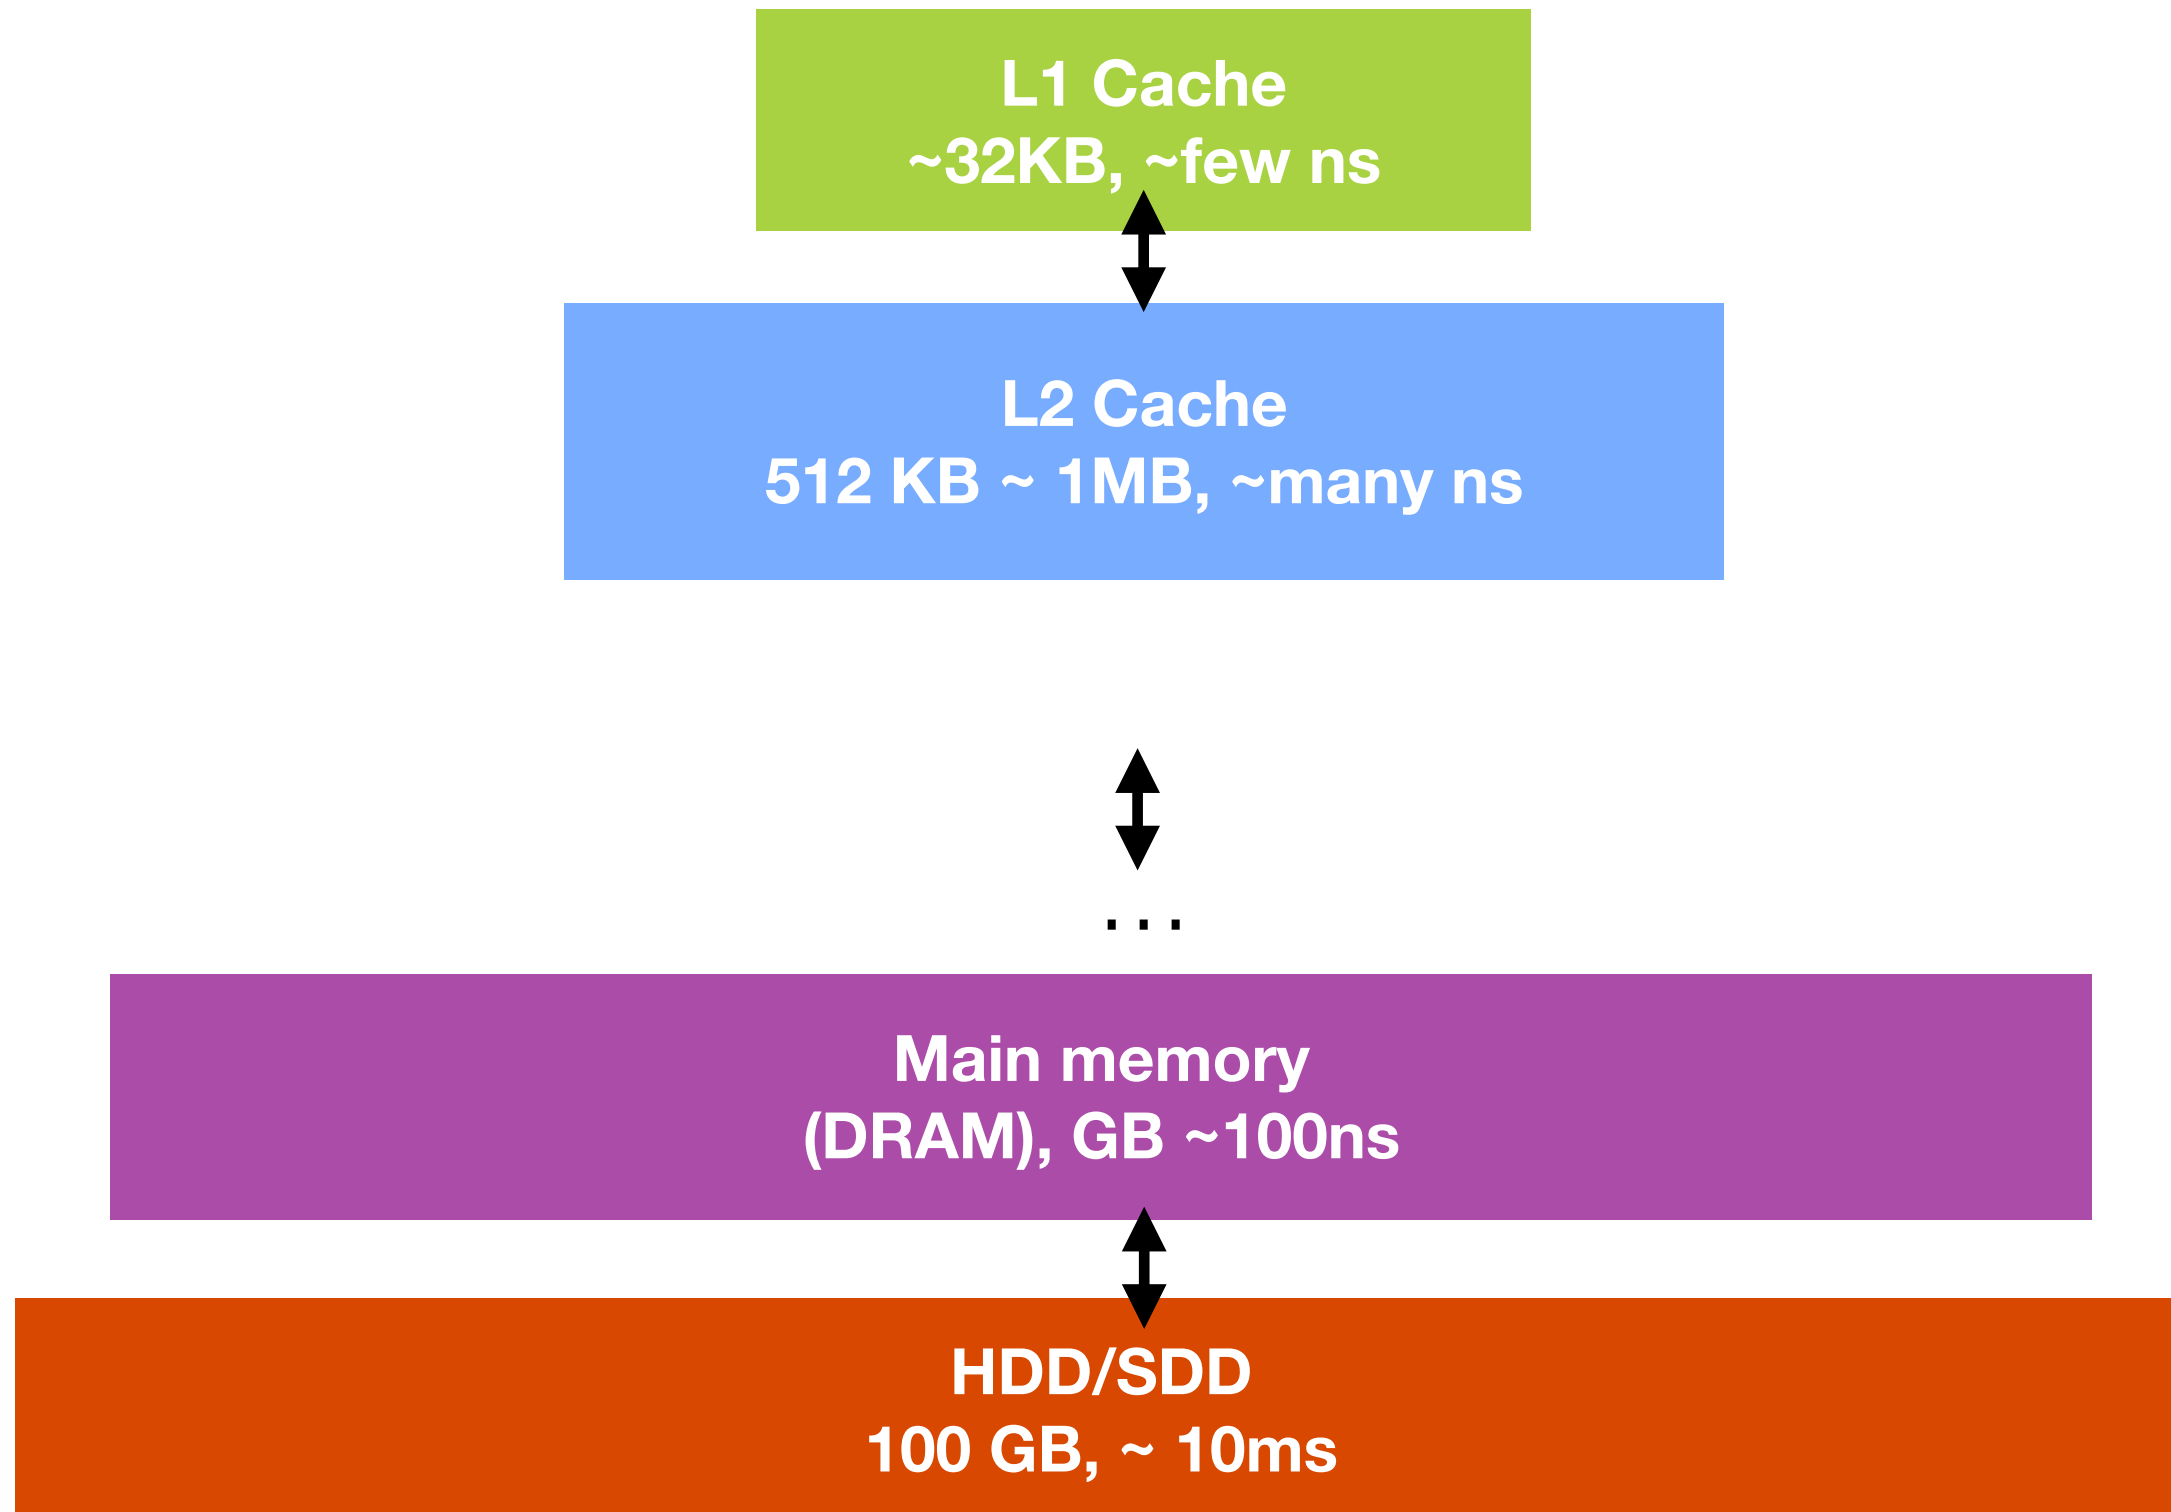
Recall: Memory Hierarchy



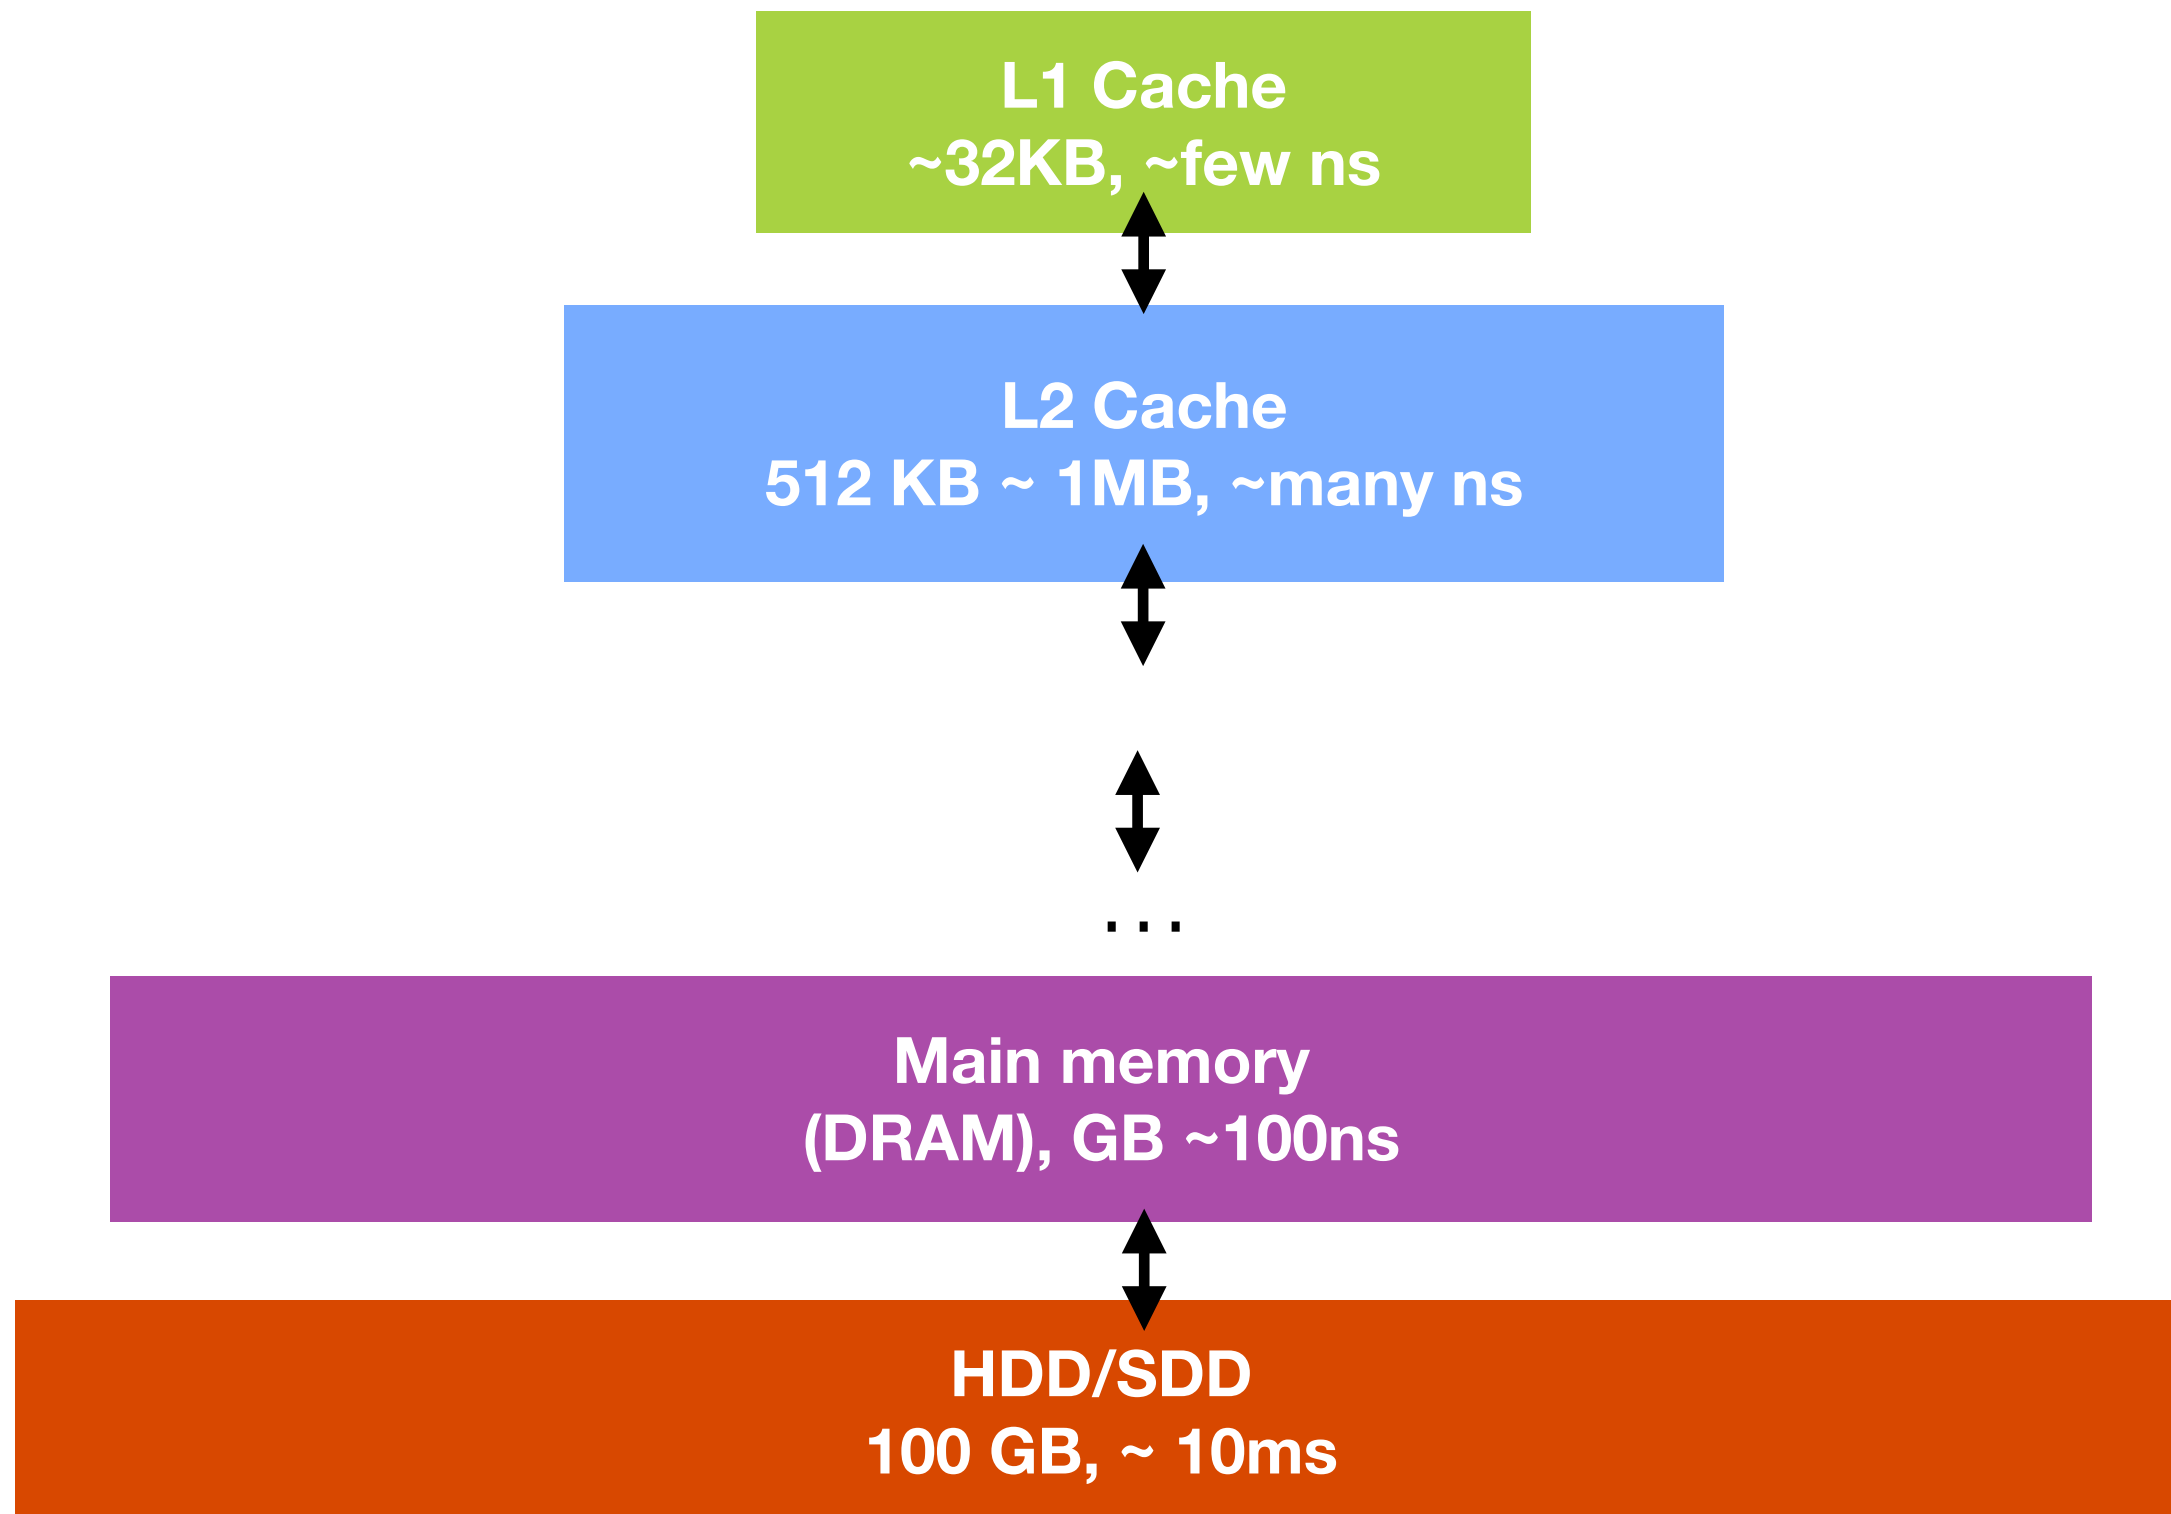
Recall: Memory Hierarchy



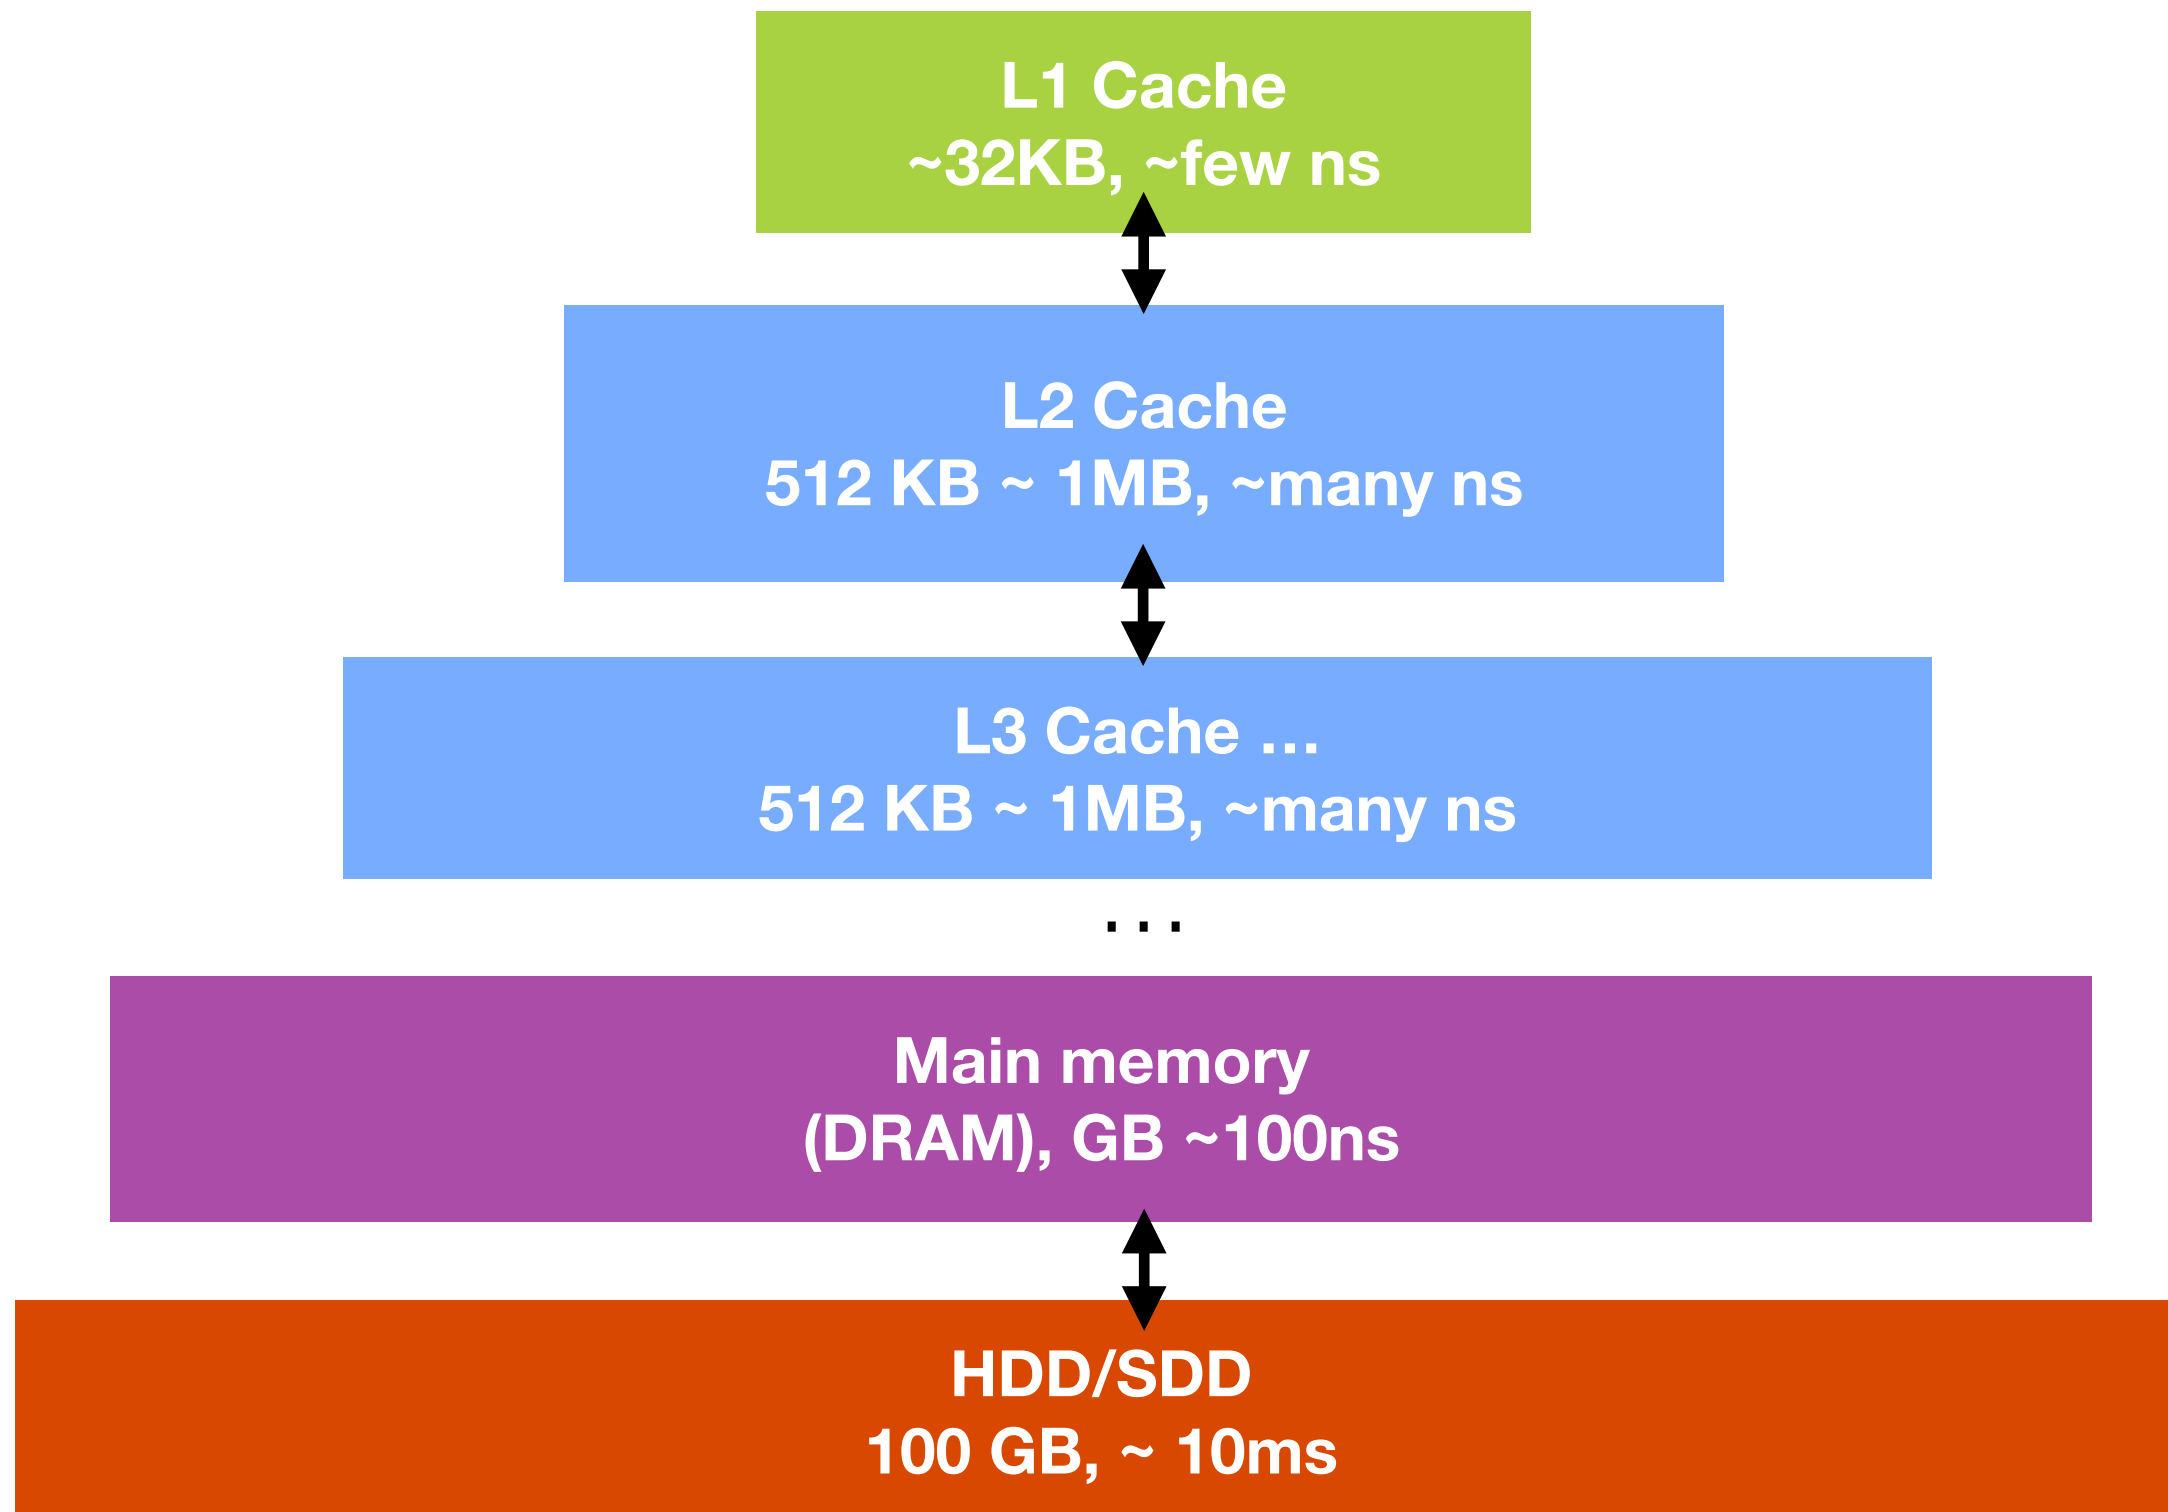
Recall: Memory Hierarchy



Recall: Memory Hierarchy



Recall: Memory Hierarchy



Recall: Cache

- **Definition:** Structure to “**store**” the recently and/or frequently used data and results to avoid high-latency operations on other structures to generate new data and results.
- This concept extends to OS: **Demand Paging**.
- **Cache basics:**
 - **Block (line):** Cache storage unit.
 - **Hit:** if in cache, the cached data is used instead of accessing the next level.
 - **Miss:** if not in cache, access the next level and stored the data inside (new cached data).

Recall: Memory Hierarchy


...

Recall: Memory Hierarchy

L1 Cache
~32KB, ~few ns


...

Recall: Memory Hierarchy

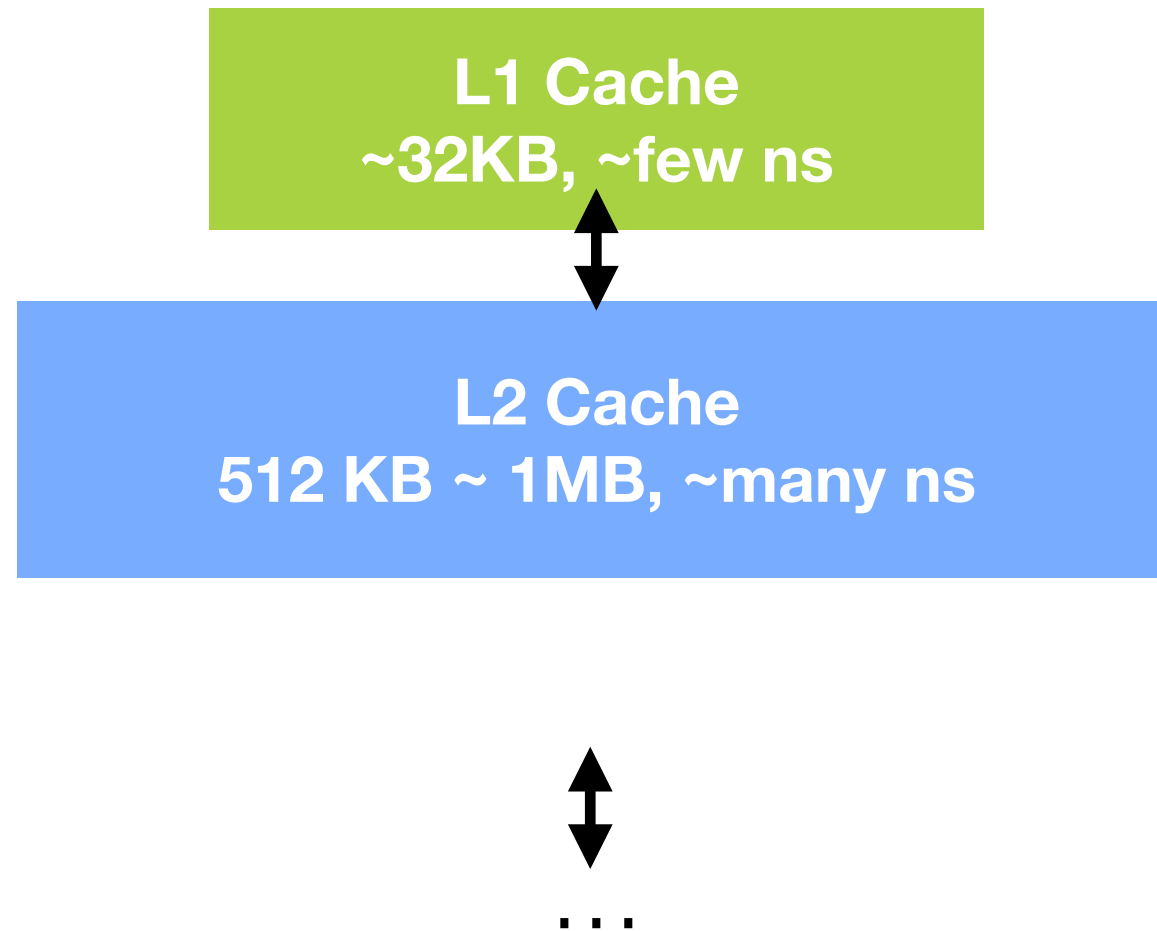


...

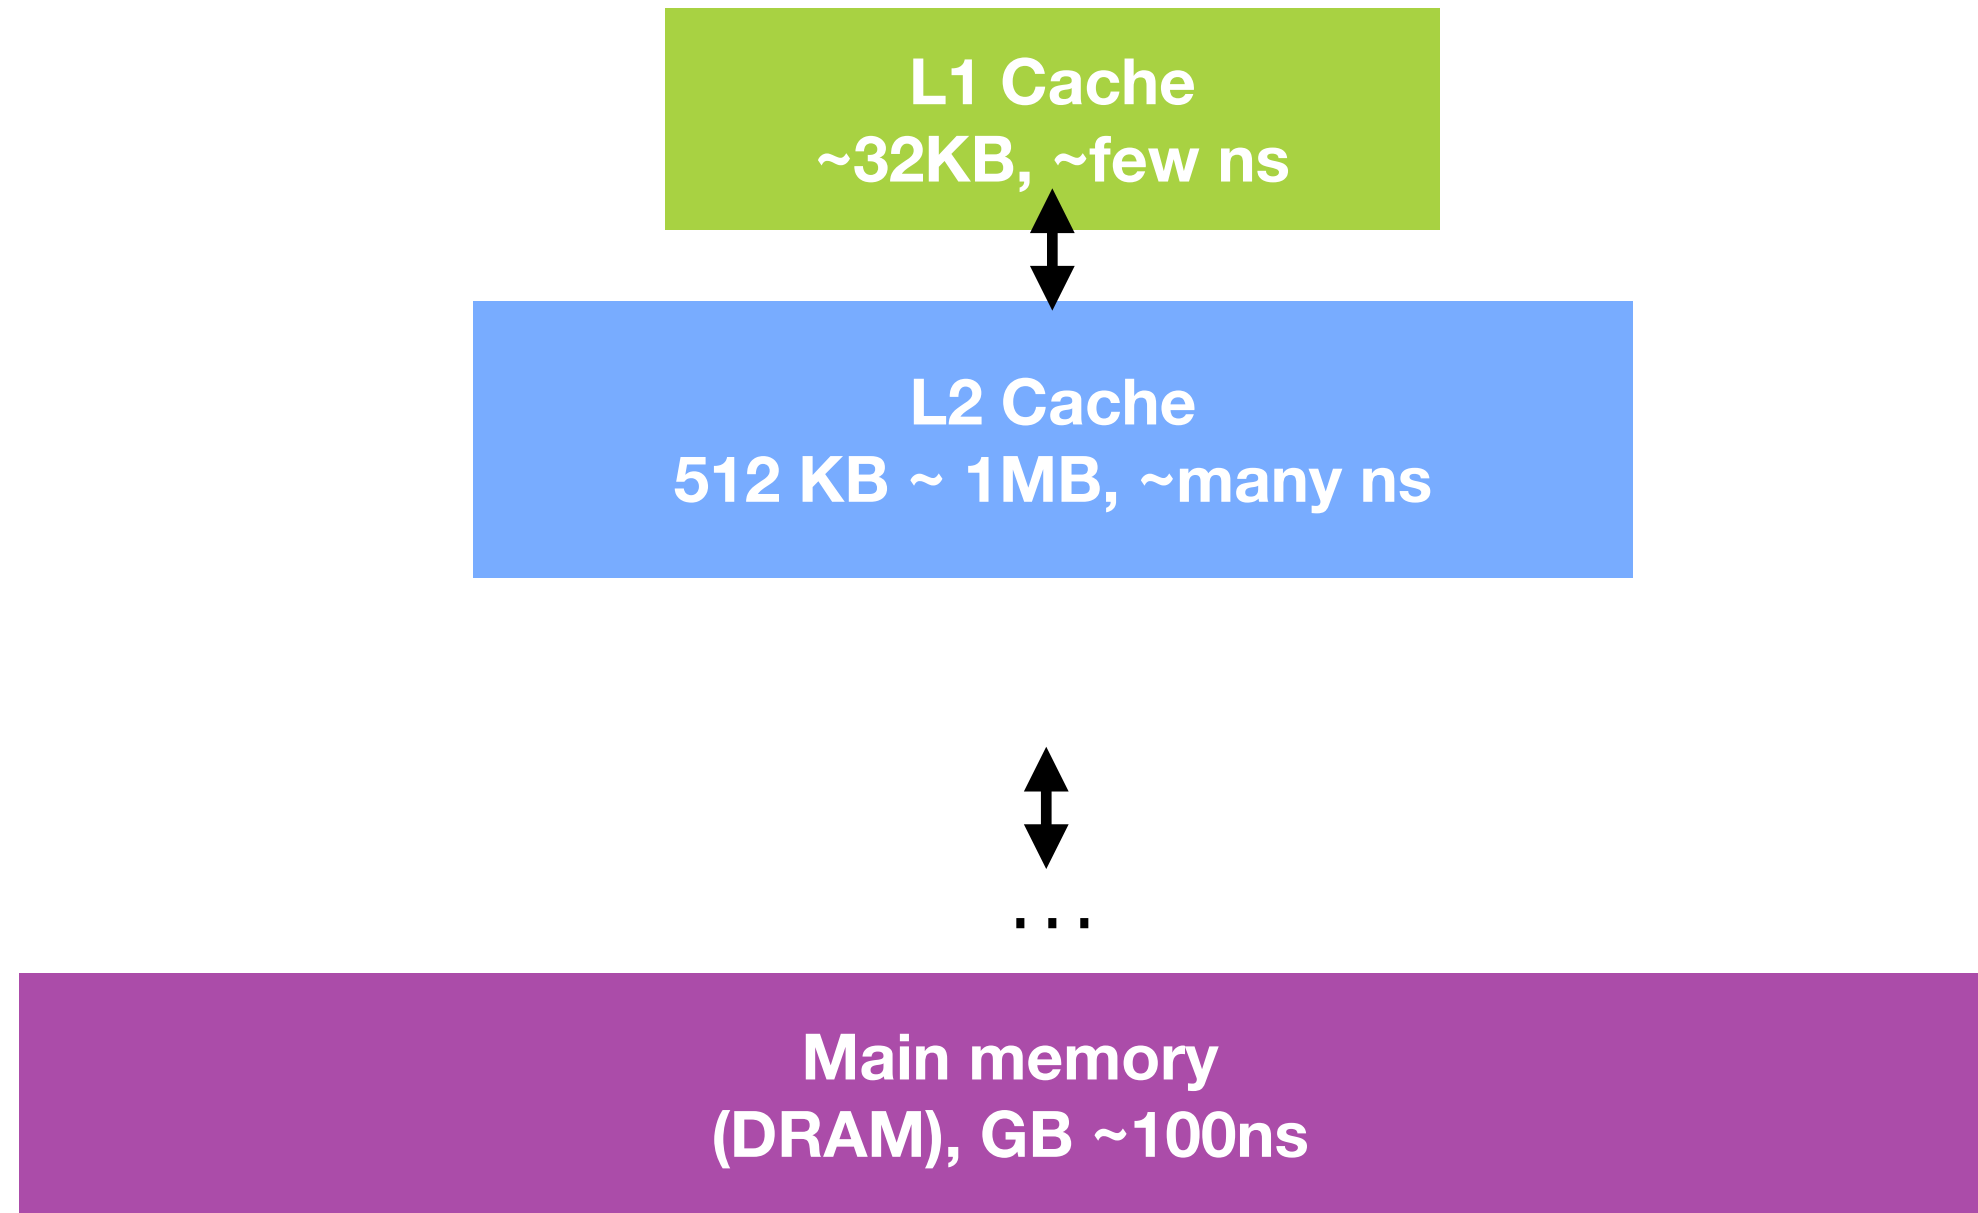
Recall: Memory Hierarchy



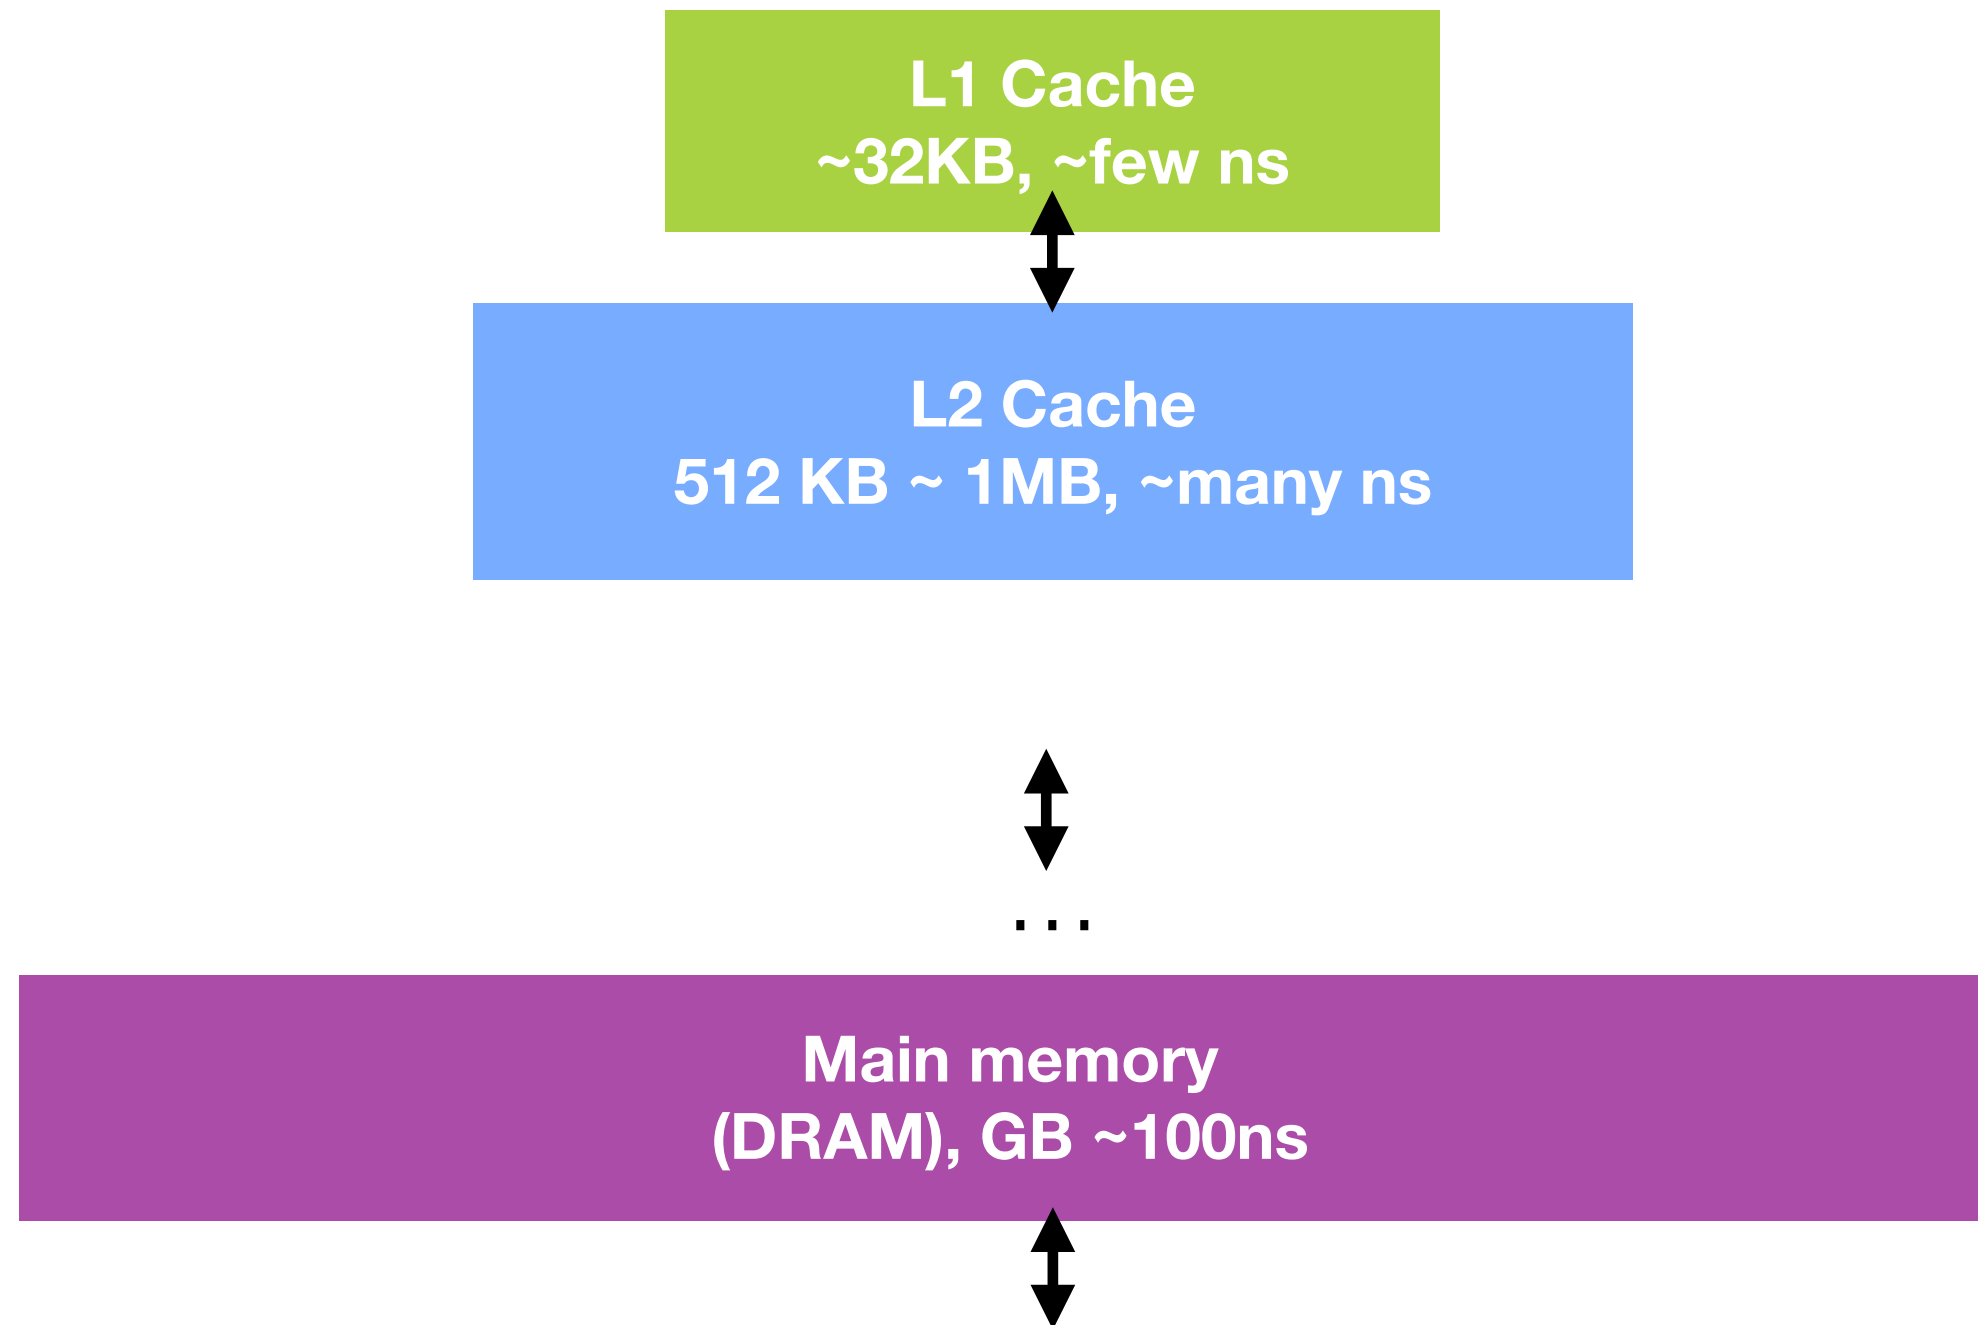
Recall: Memory Hierarchy



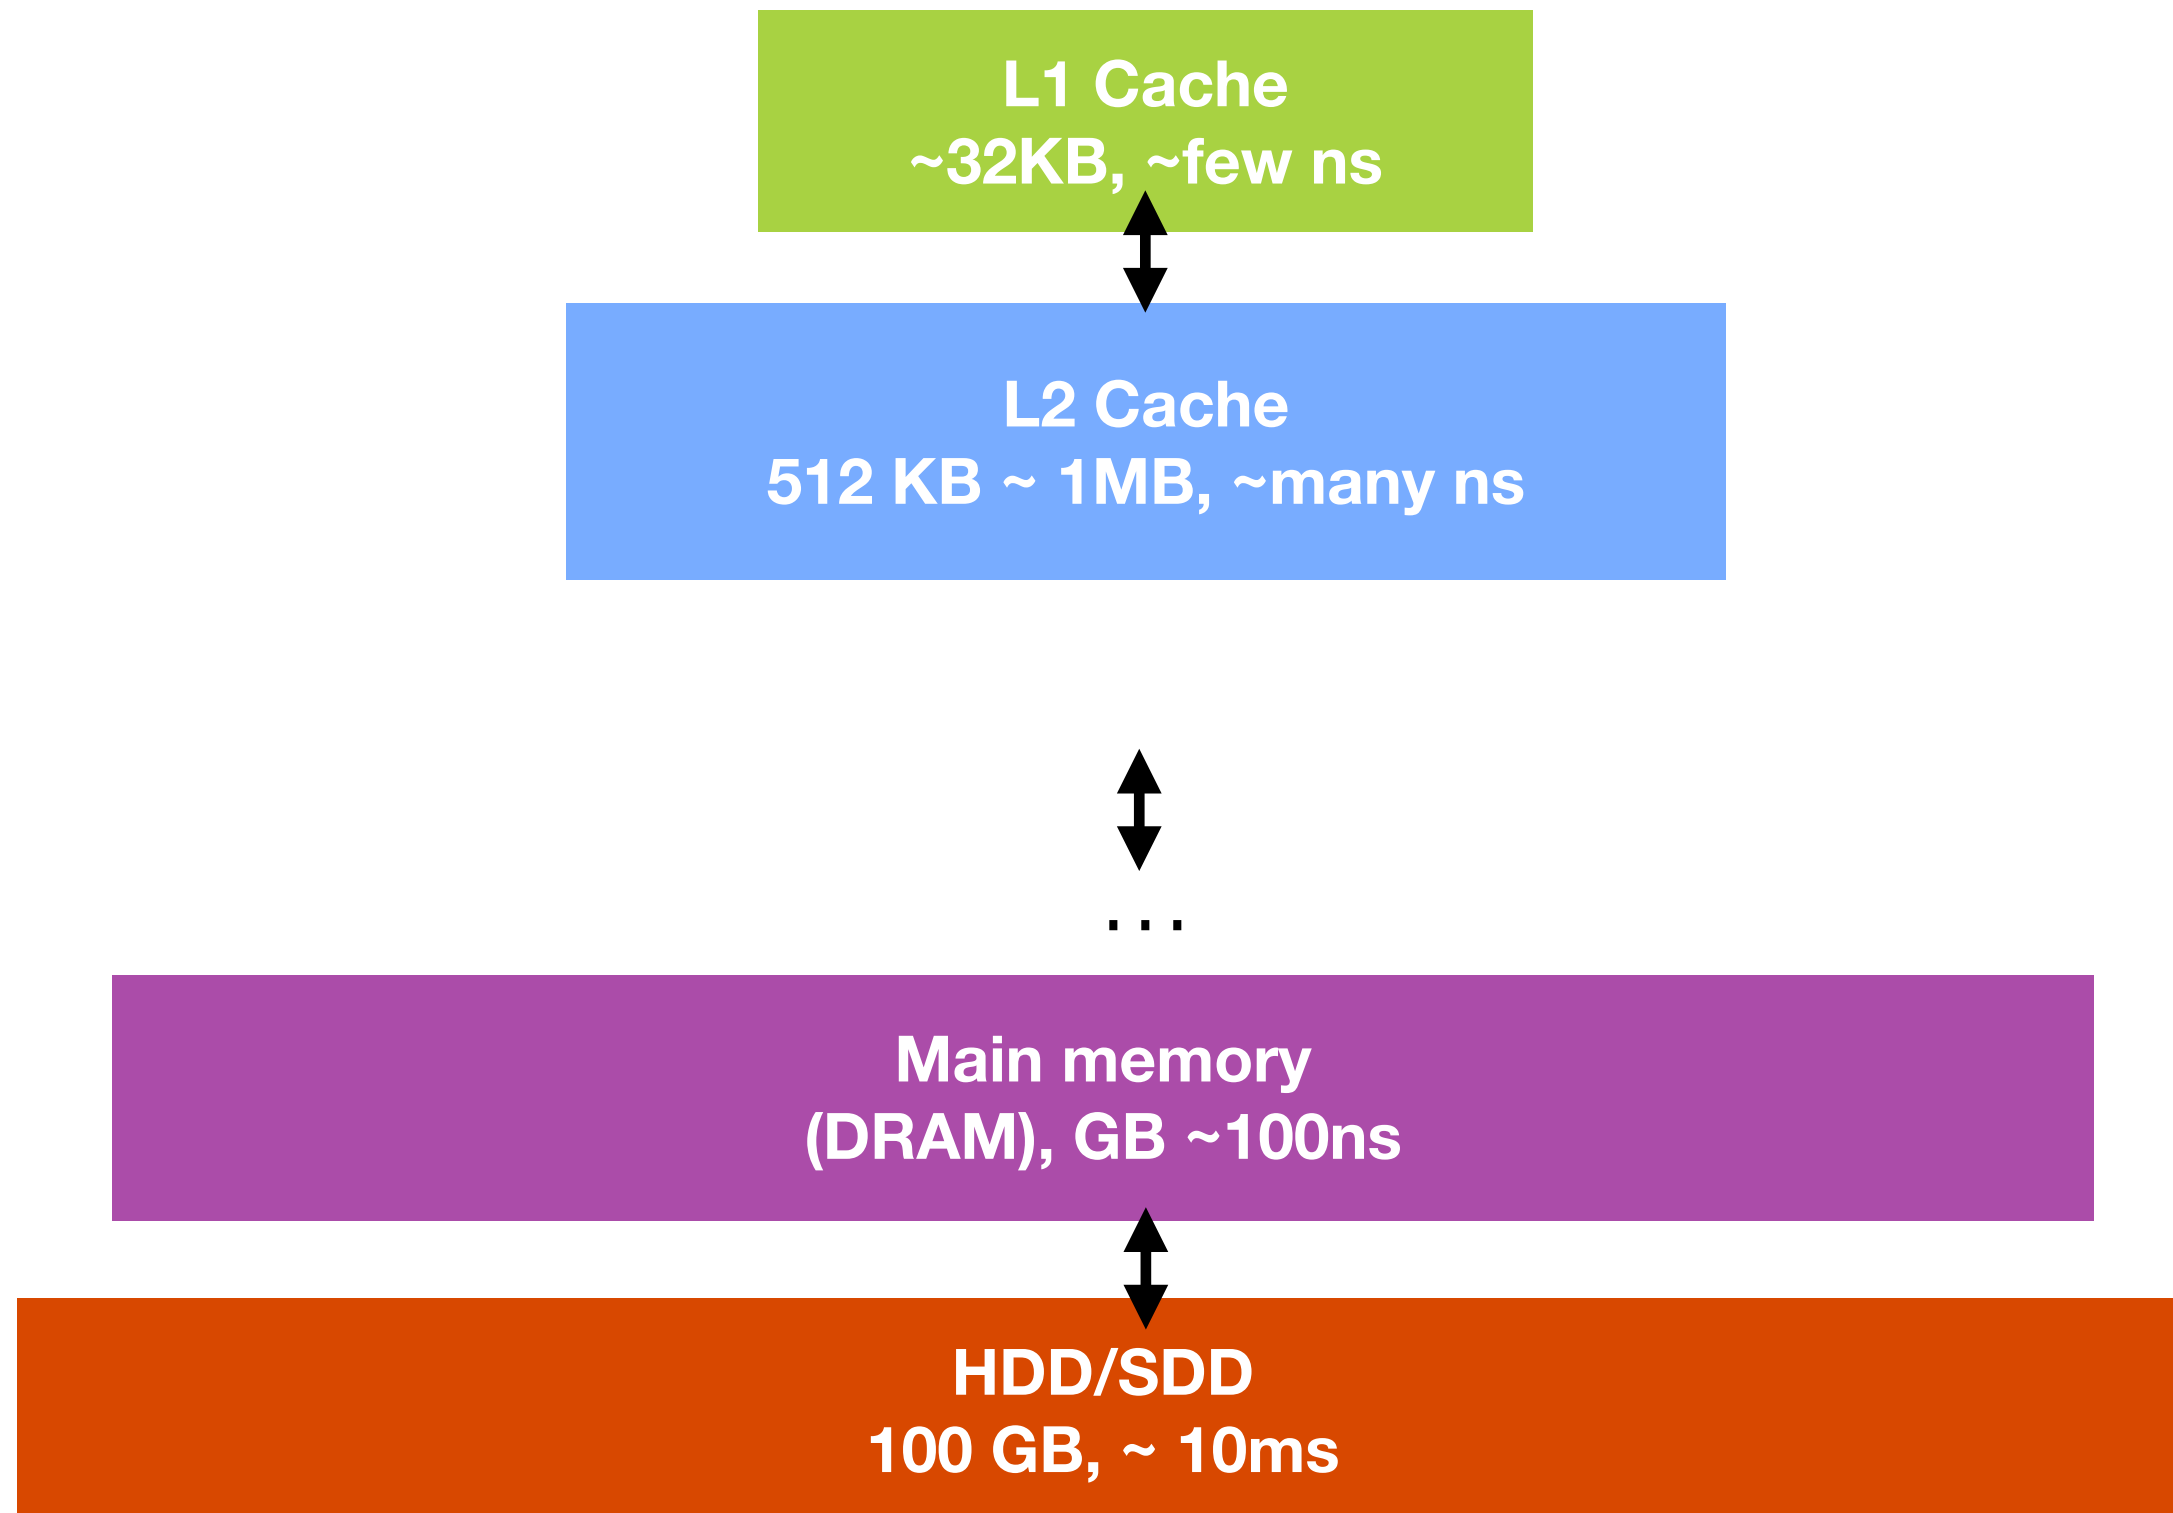
Recall: Memory Hierarchy



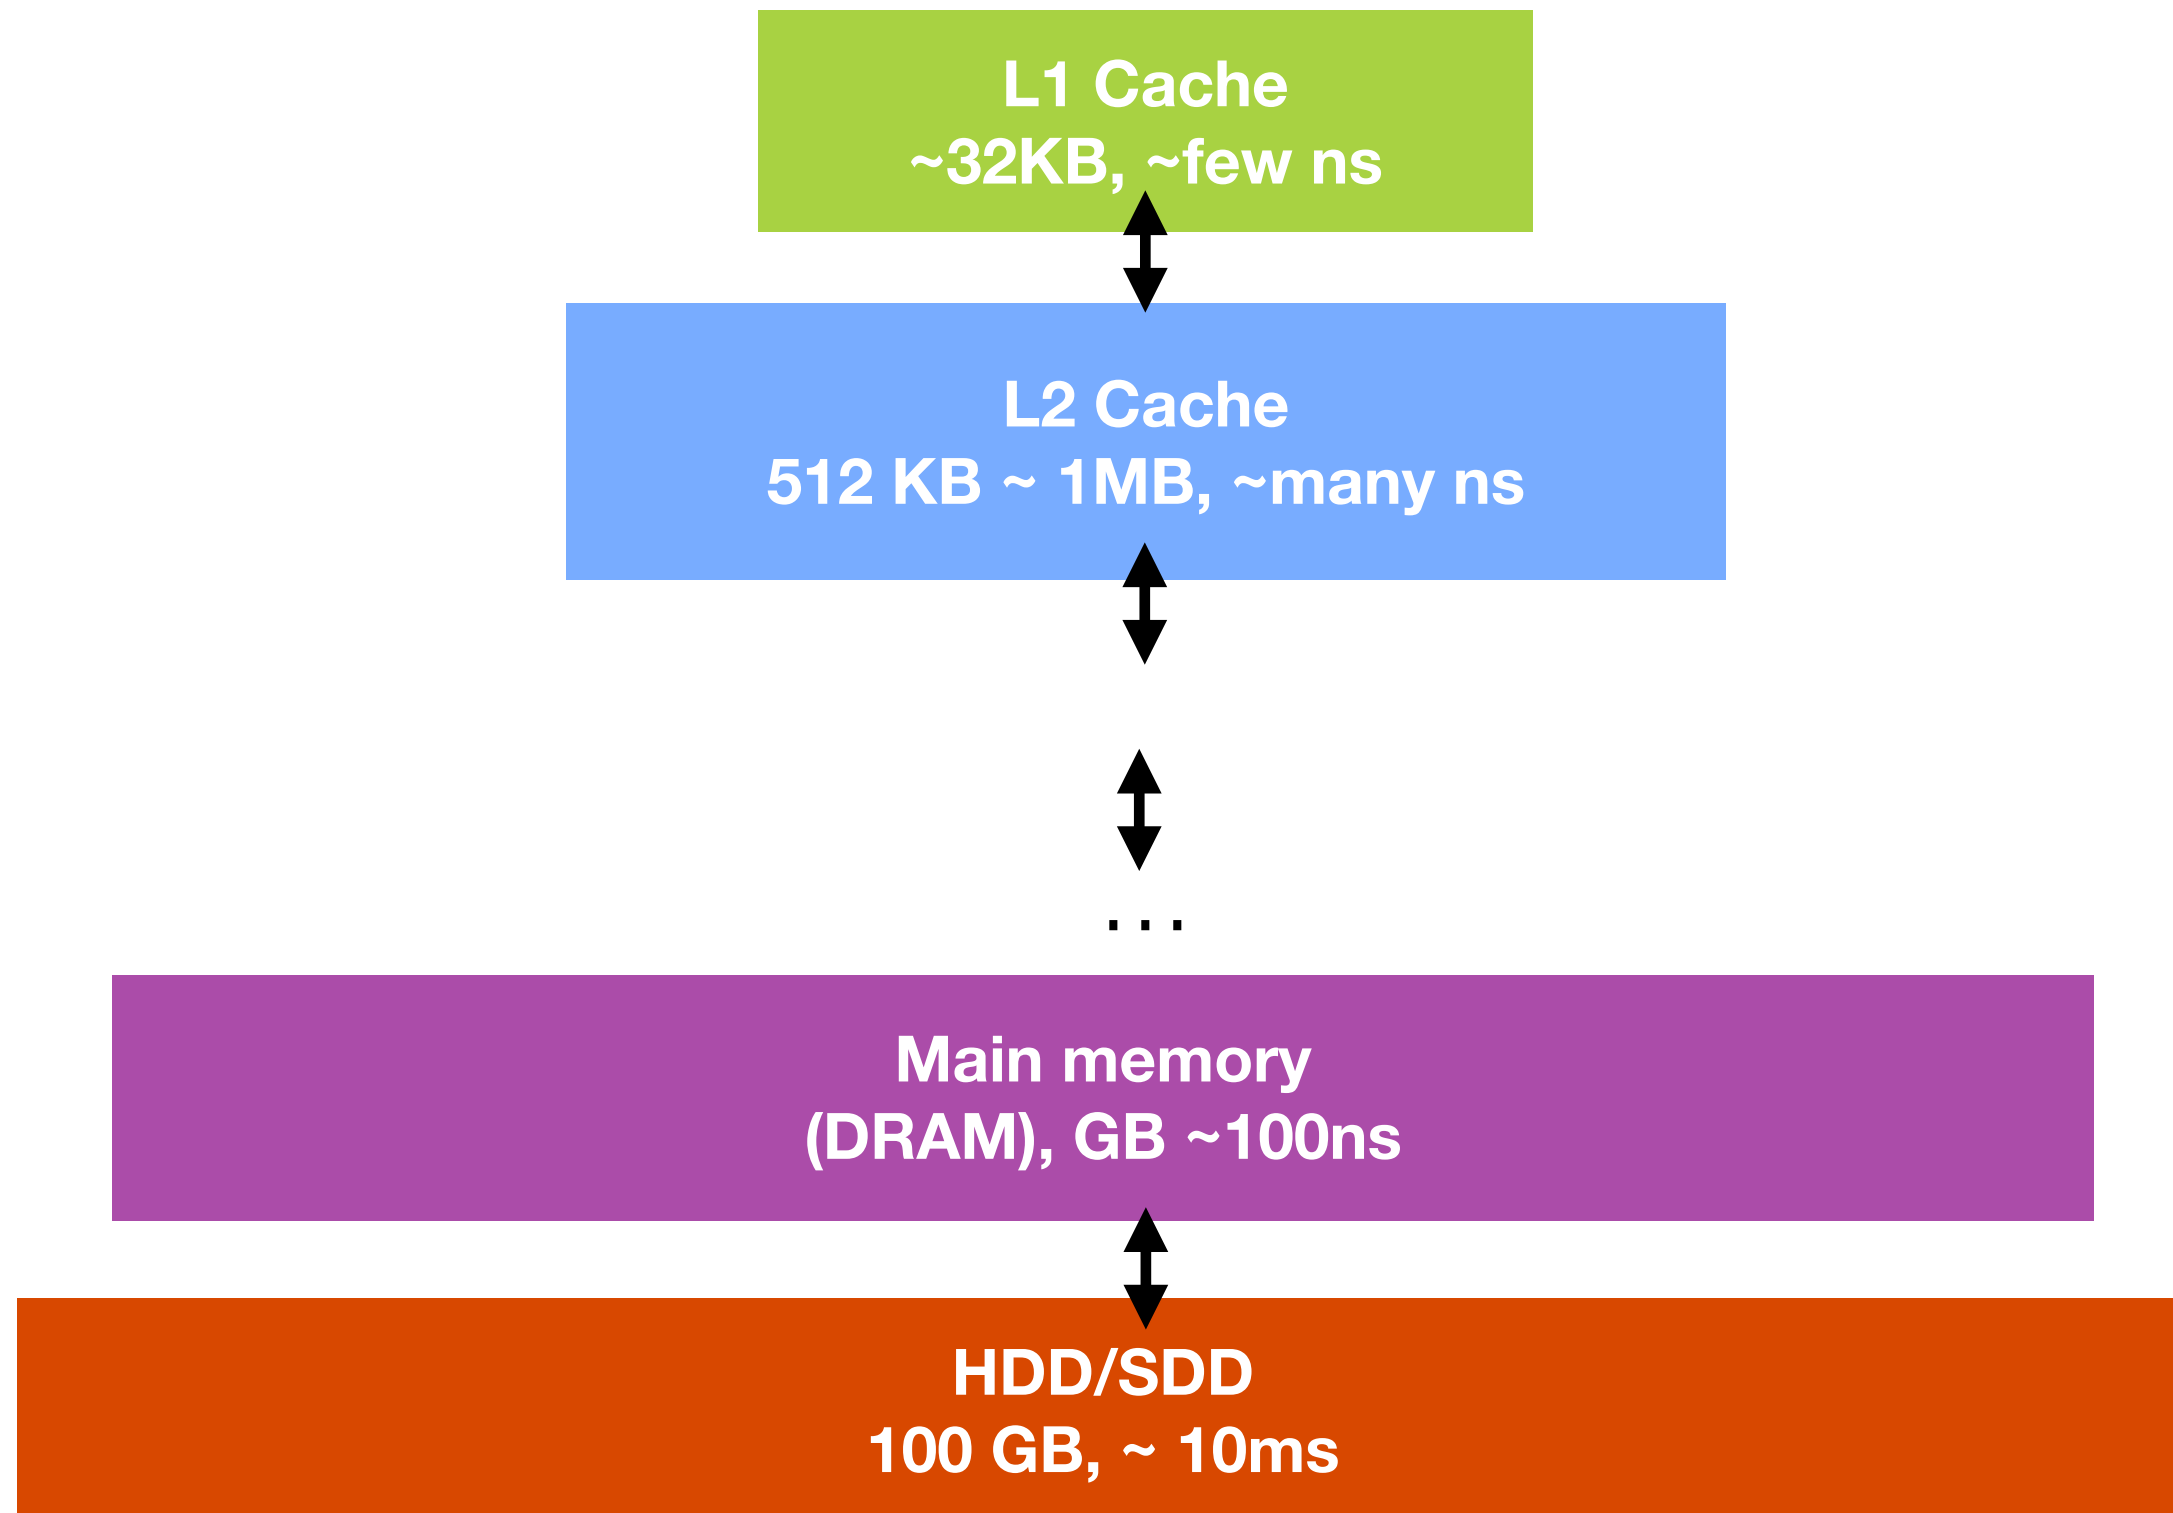
Recall: Memory Hierarchy



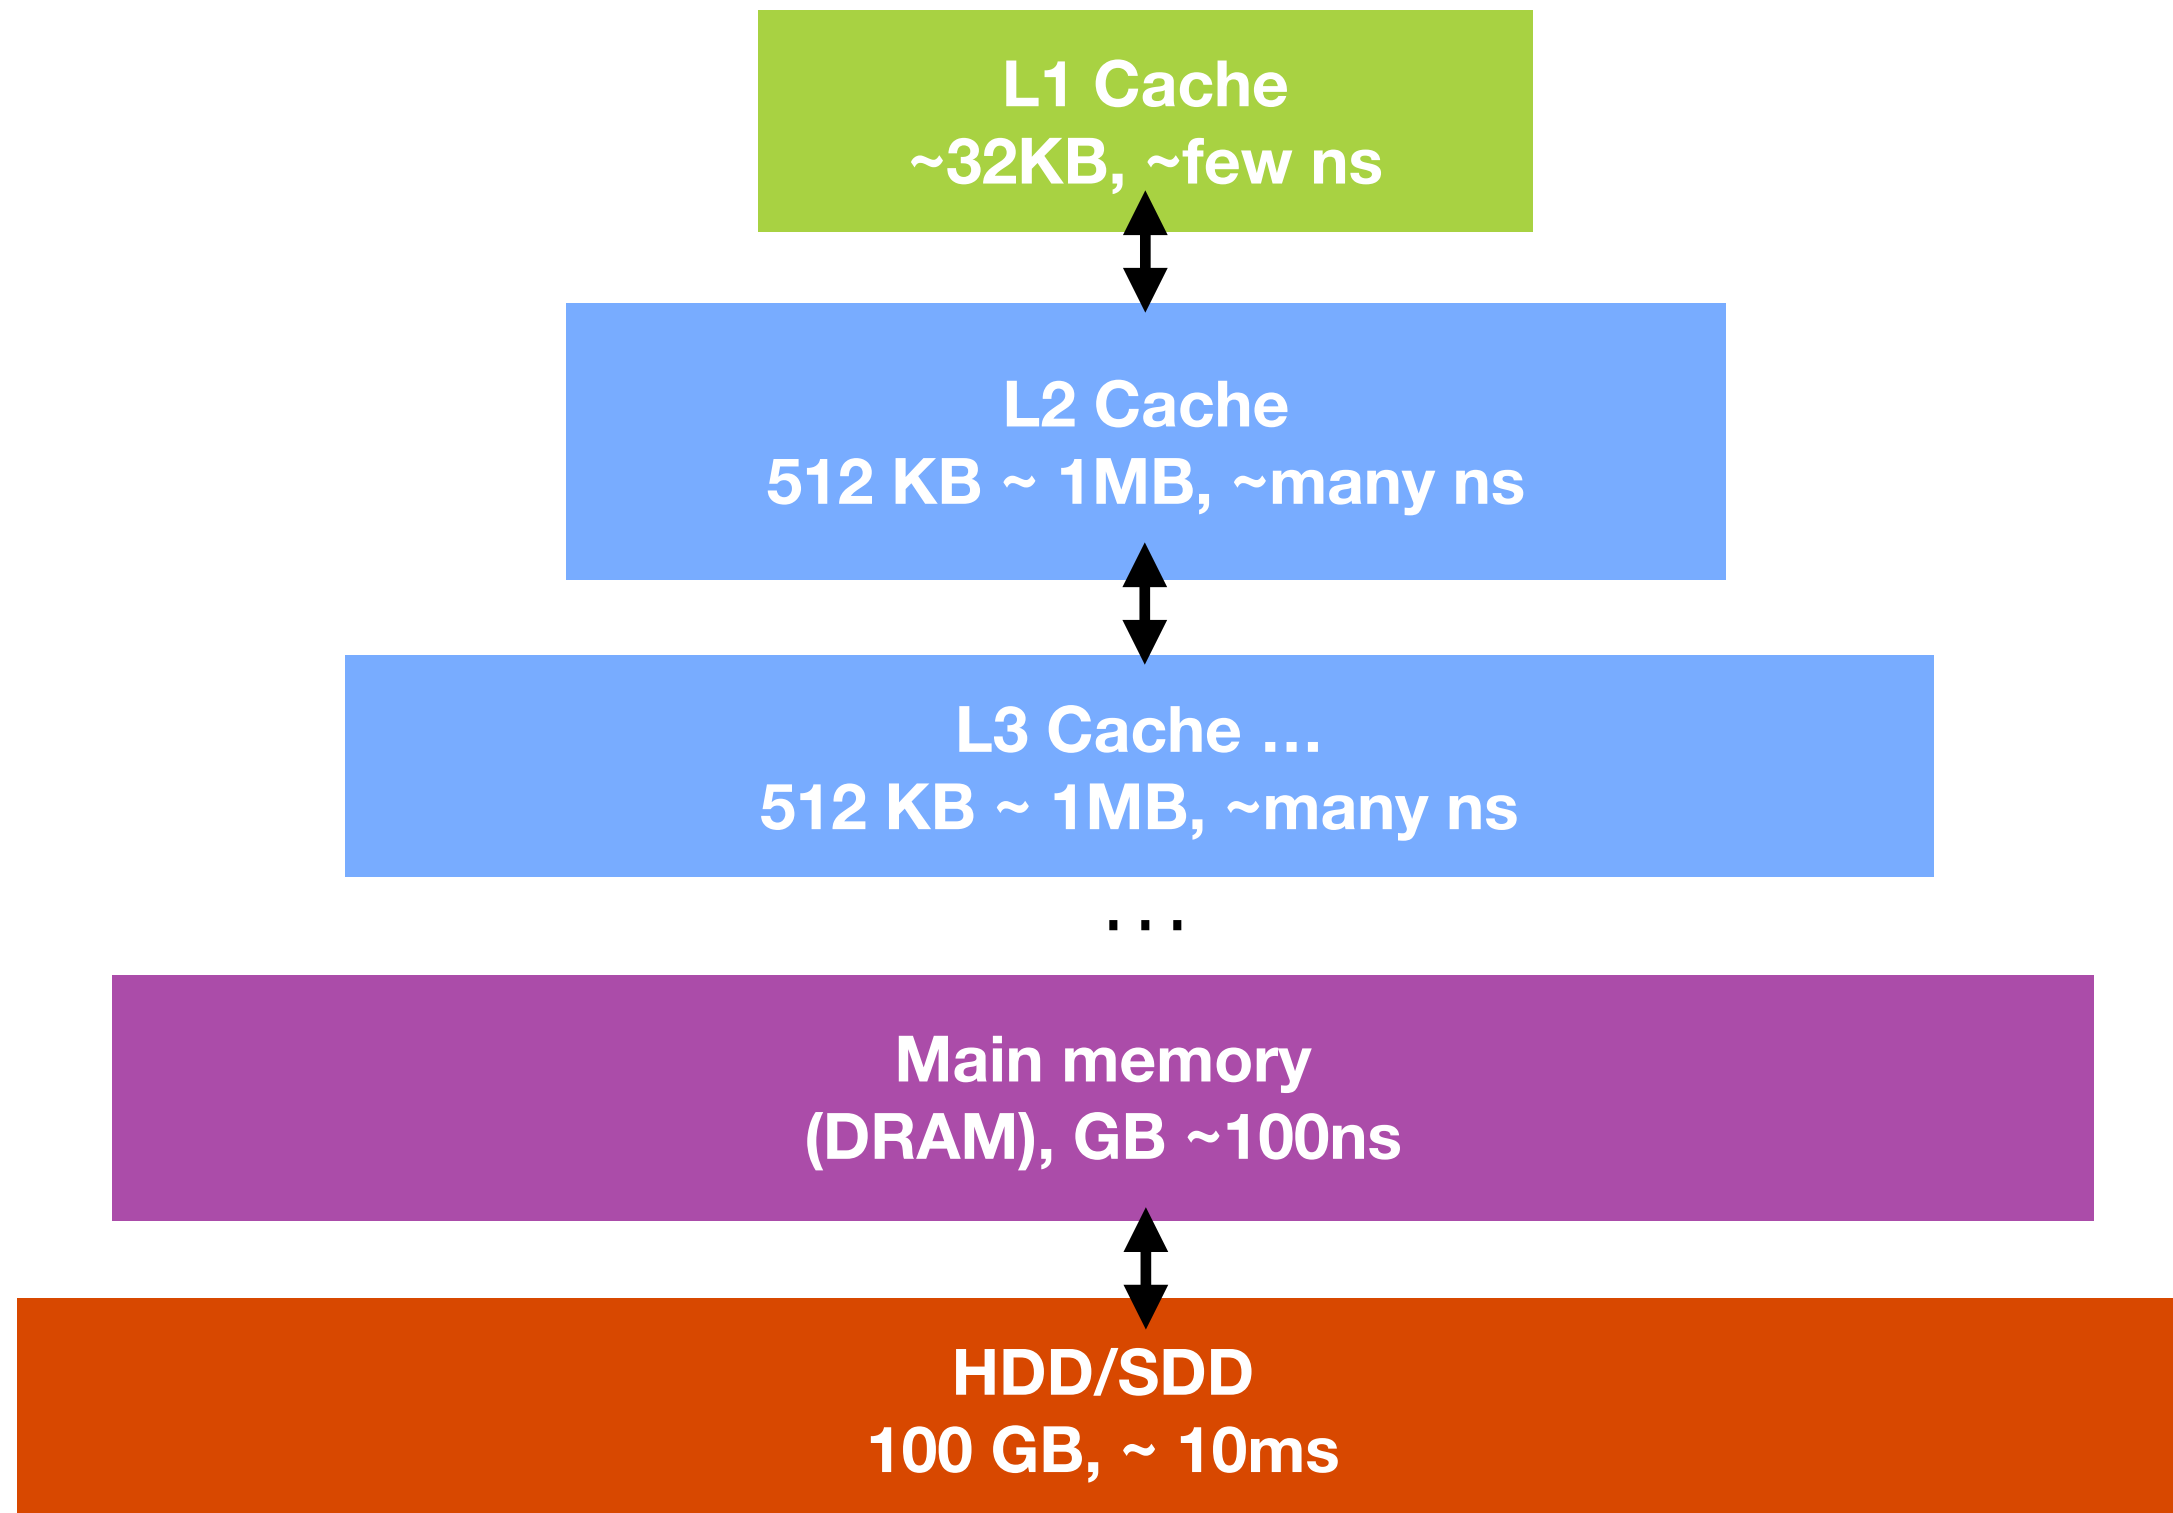
Recall: Memory Hierarchy



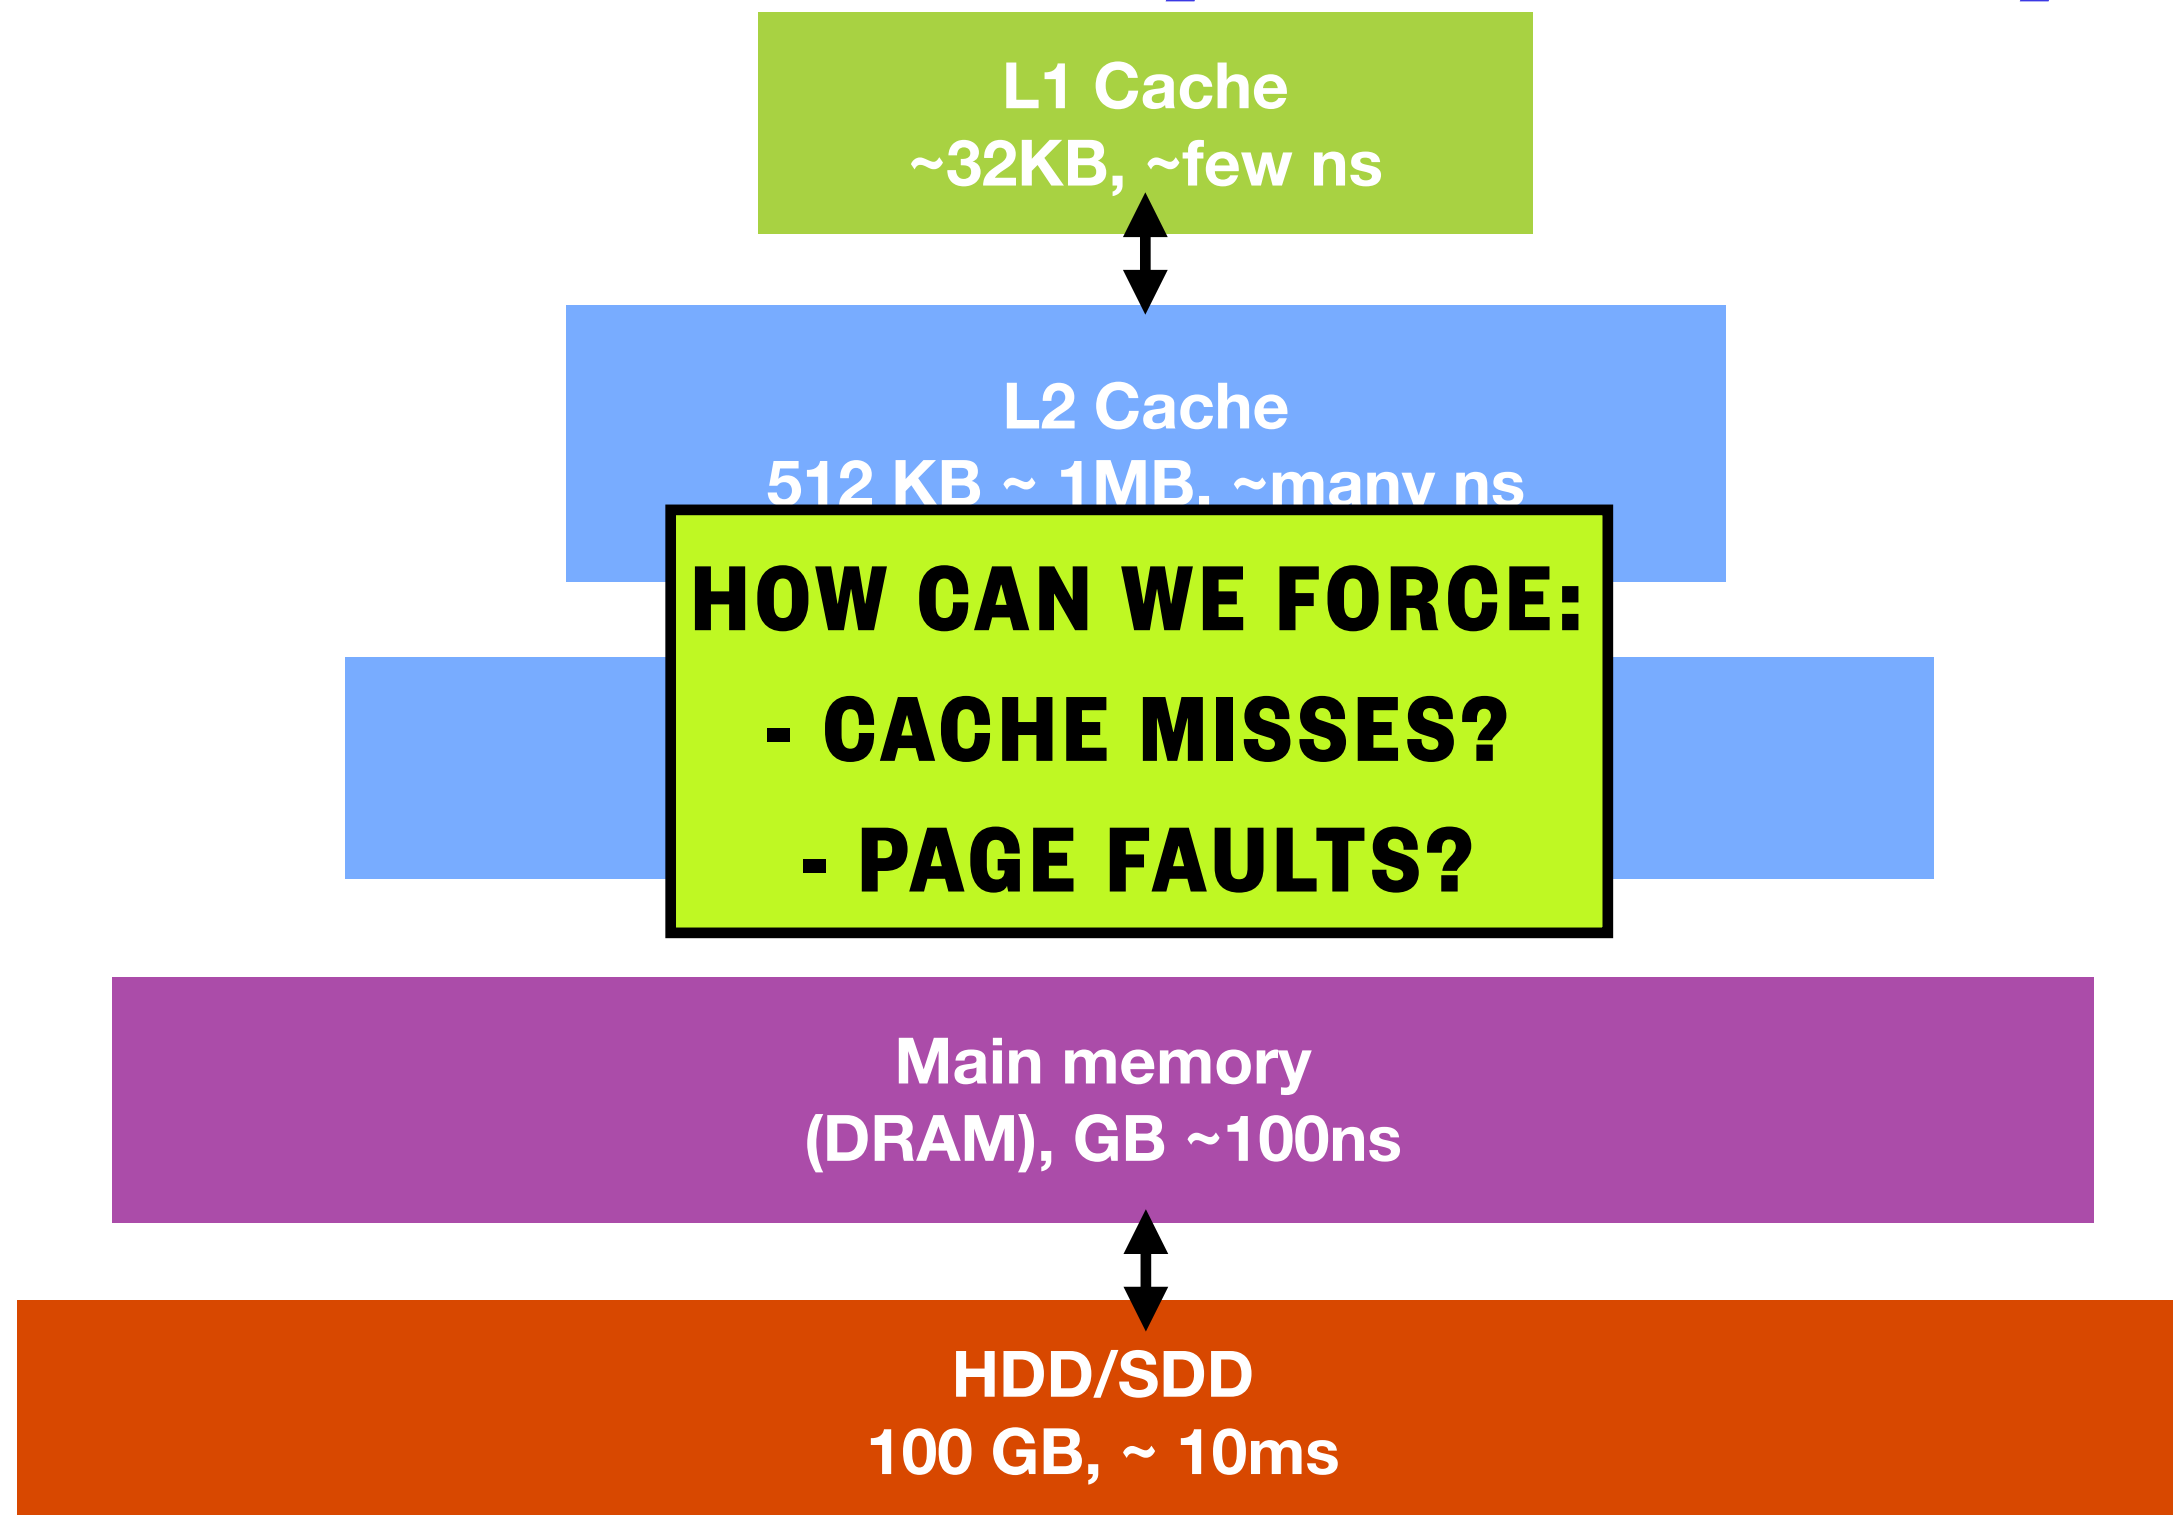
Recall: Memory Hierarchy



Recall: Memory Hierarchy




Recall: Memory Hierarchy



Exercise 1:

1. Identify the architecture of your system: cache levels, cache size, block size, frequency, dram size, tlb levels. **Summarize in a table.**
2. **Implement a program Pfaults** to force misses and page faults:
 1. Create an array that exceeds the DRAM size. **Notice:** are we using virtual memory ;) ?
 2. Perform random accesses over the array in memory (reads/loads).
 3. Control the execution time of your program (define **X** minutes).
3. Obtain all the data from cache (misses), TLB (misses), page faults, free memory using **perf** of the (total execution time **2X** minutes) :
 1. **First X minutes:** in your normal system state, while you are working and not executing the Pfaults program.
 2. **Last X minutes:** while executing the Pfaults program.
 3. Parse the perf output and plot your results. **Summarize** in a table.
4. Explain your plot. Justify.

Plot example



TODO: Check with perf:

```
# Various CPU level 1 data cache statistics for the  
specified command:
```

```
perf stat -e L1-dcache-loads,L1-dcache-load-misses,L1-  
dcache-stores command
```

```
# Various CPU data TLB statistics for the specified  
command:
```

```
perf stat -e dTLB-loads,dTLB-load-misses,dTLB-prefetch-  
misses command
```

```
perf stat -e cpu-clock ./programa
```

```
perf stat -e cpu-clock,faults ./programa
```

```
perf report --stdio --sort comm,dso
```


IOstat

- Install IOstat:

```
sudo apt-get install sysstat
```

- Examples

```
$iostat
```

```
$iostat -d 5 3
```

```
Linux 3.19.0-25-generic (Ubuntu-PC)      Saturday 16 December 2017  
_x86_64_ (4 CPU)
```

Device:	tps	kB_read/s	kB_wrtn/s	kB_read	kB_wrtn
sda	11.77	340.71	98.95	771022	223928

Device:	tps	kB_read/s	kB_wrtn/s	kB_read	kB_wrtn
sda	2.00	0.00	8.00	0	40

Device:	tps	kB_read/s	kB_wrtn/s	kB_read	kB_wrtn
sda	0.60	0.00	3.20	0	16

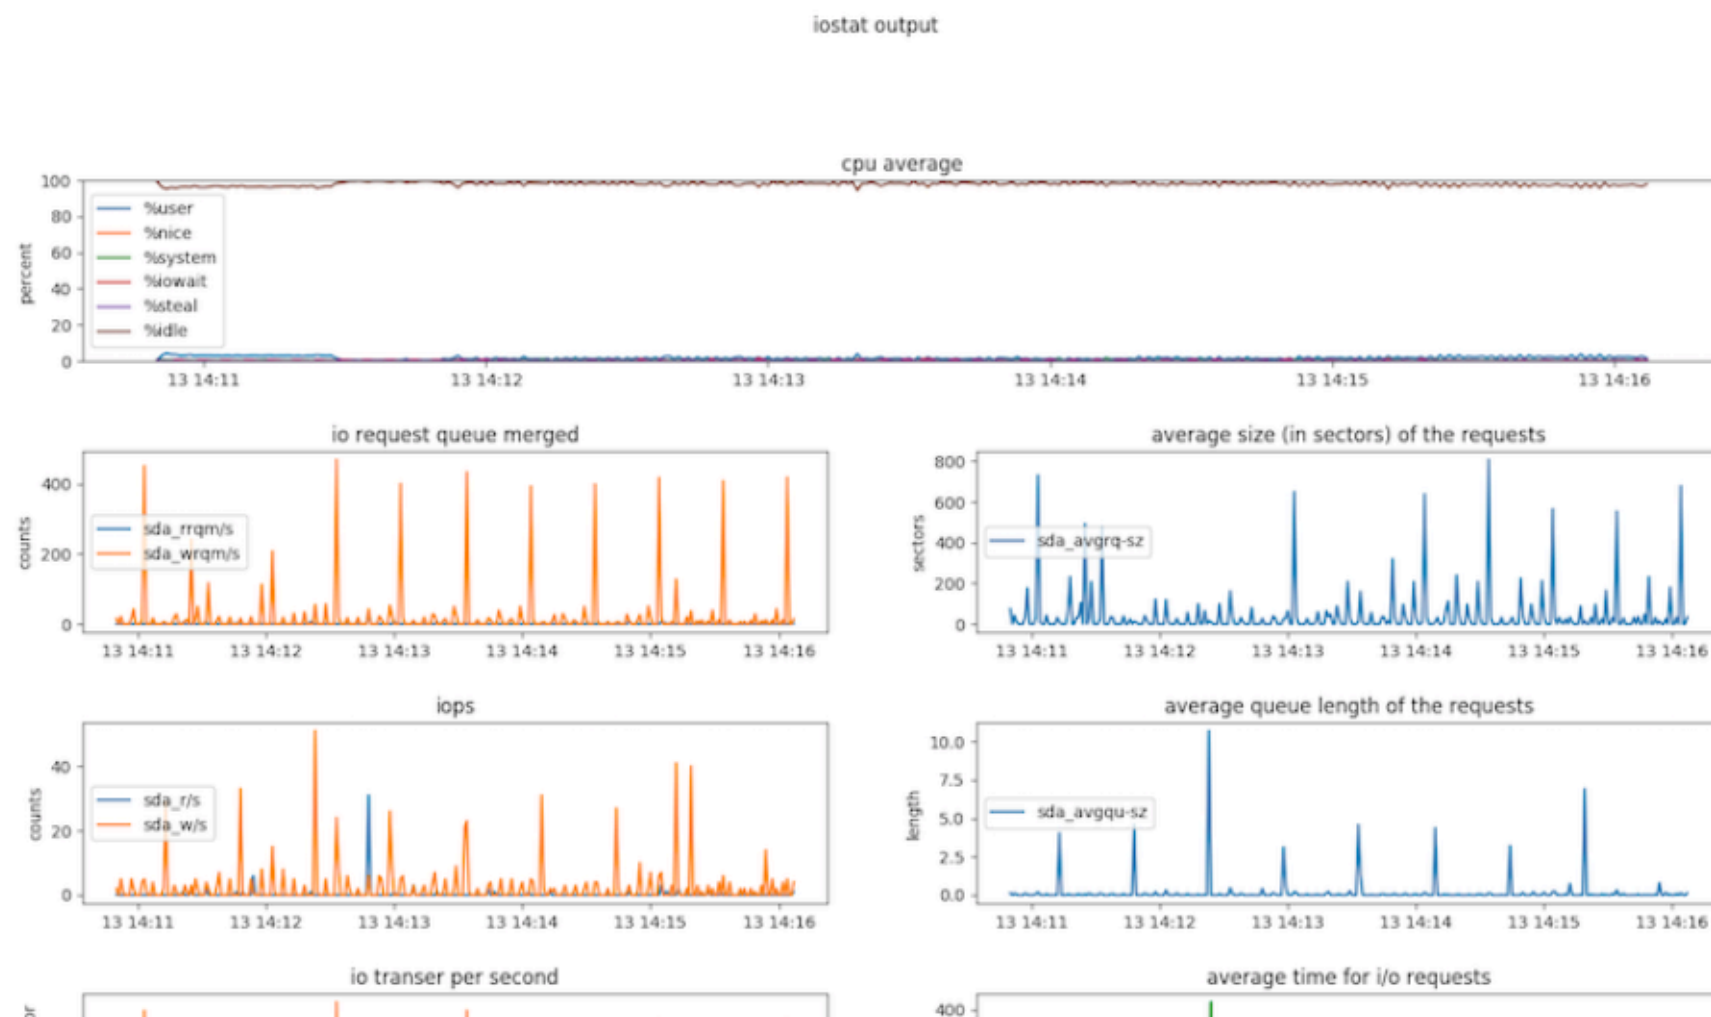
Iostat tool

- Python library for parsing Iostat results.

[HTTPS://PYPI.ORG/PROJECT/IOSTAT-TOOL/](https://pypi.org/project/iostat-tool/)

How to use

This is sample image rendered by matplotlib.



Io streams with Python

[HTTPS://DOCS.PYTHON.ORG/3/LIBRARY/IO.HTML](https://docs.python.org/3/library/io.html)

Python » English » 3.8.3 » Documentation » The Python Standard Library » Generic Operating System Services »

previous | next

Quick

io — Core tools for working with streams

Source code: [Lib/io.py](#)

Overview ¶

The `io` module provides Python's main facilities for dealing with various types of I/O. There are three main types of I/O: *text I/O*, *binary I/O* and *raw I/O*. These are generic categories, and various backing stores can be used for each of them. A concrete object belonging to any of these categories is called a [file object](#). Other common terms are *stream* and *file-like object*.

Independent of its category, each concrete stream object will also have various capabilities: it can be read-only, write-only, or read-write. It can also allow arbitrary random access (seeking forwards or backwards to any location), or only sequential access (for example in the case of a socket or pipe).

All streams are careful about the type of data you give to them. For example giving a `str` object to the `write()` method of a binary stream will raise a `TypeError`. So will giving a `bytes` object to the `write()` method of a text stream.

Changed in version 3.3: Operations that used to raise `IOError` now raise `OSError`, since `IOError` is now an alias of `OSError`.

Table of Contents

io — Core tools for working with streams

- Overview
 - Text I/O
 - Binary I/O
 - Raw I/O
- High-level Module Interface
 - In-memory streams
- Class hierarchy
 - I/O Base Classes
 - Raw File I/O
 - Buffered Streams
 - Text I/O
- Performance
 - Binary I/O
 - Text I/O
 - Multi-threading
 - Reentrancy

Previous topic

os — Miscellaneous operating system interfaces

Exercise 2

- Perform I/O (read and write) operations.
 1. Implement a program with the following functions:
 - A. Random reads size 200MB.
 - B. Random writes size 200MB.
 - C. Sequential reads size 200 MB.
 - D. Sequential writes size 200 MB.
 - E. Random writes and reads with variable size (max 500MB).
 - F. Define an amount of X minutes for each function call (max 2min).
 2. Generate an output file using iostat during your program execution.
 - **Notice: The iostat data can reach GB file size.**
 3. Process and plot your data using the Python library for IOstat.
 - Explain your results. **Justify.**